



# About CloudThat

**9+**

Years in Business

**350k+**

Professional Trained

**100+**

Corporates Trained

**100+**

Projects Delivered

**350+**

Cloud Certifications

Leader in Training and Consulting on Cloud, Security, AI/ML, IoT and DevOps

Trained over 350k+ professionals across technologies and geographies

Proven track record of training delivery for all stages of employee lifecycle

Strong team of 125+ certified cloud experts with industry consulting experience

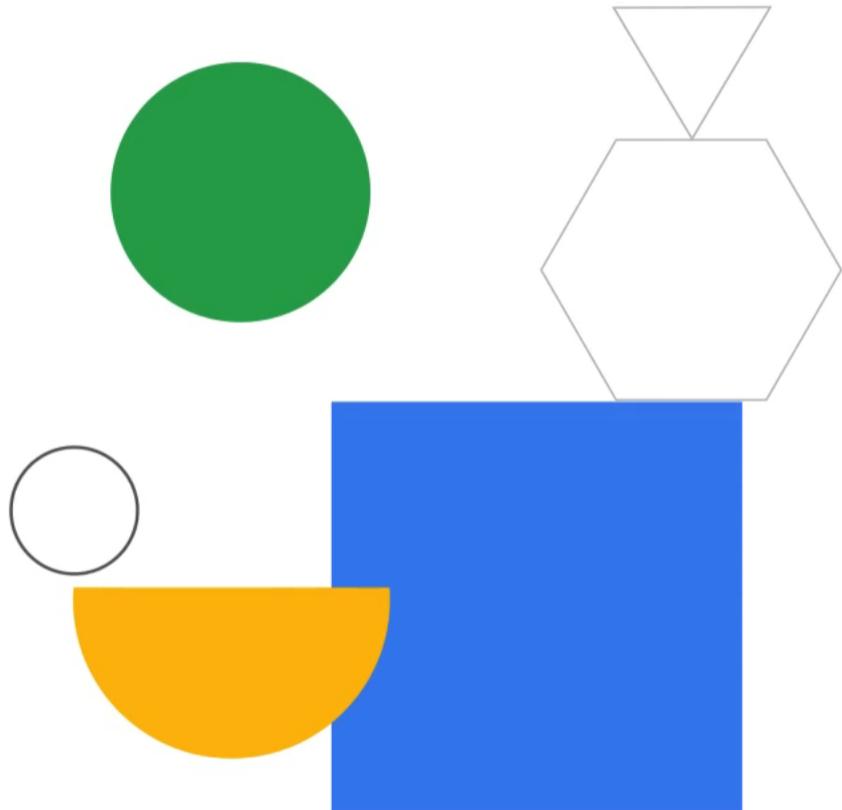
Robust consulting division brings industry prospective to training delivery



2020 Partner of the Year Finalist  
Learning Award

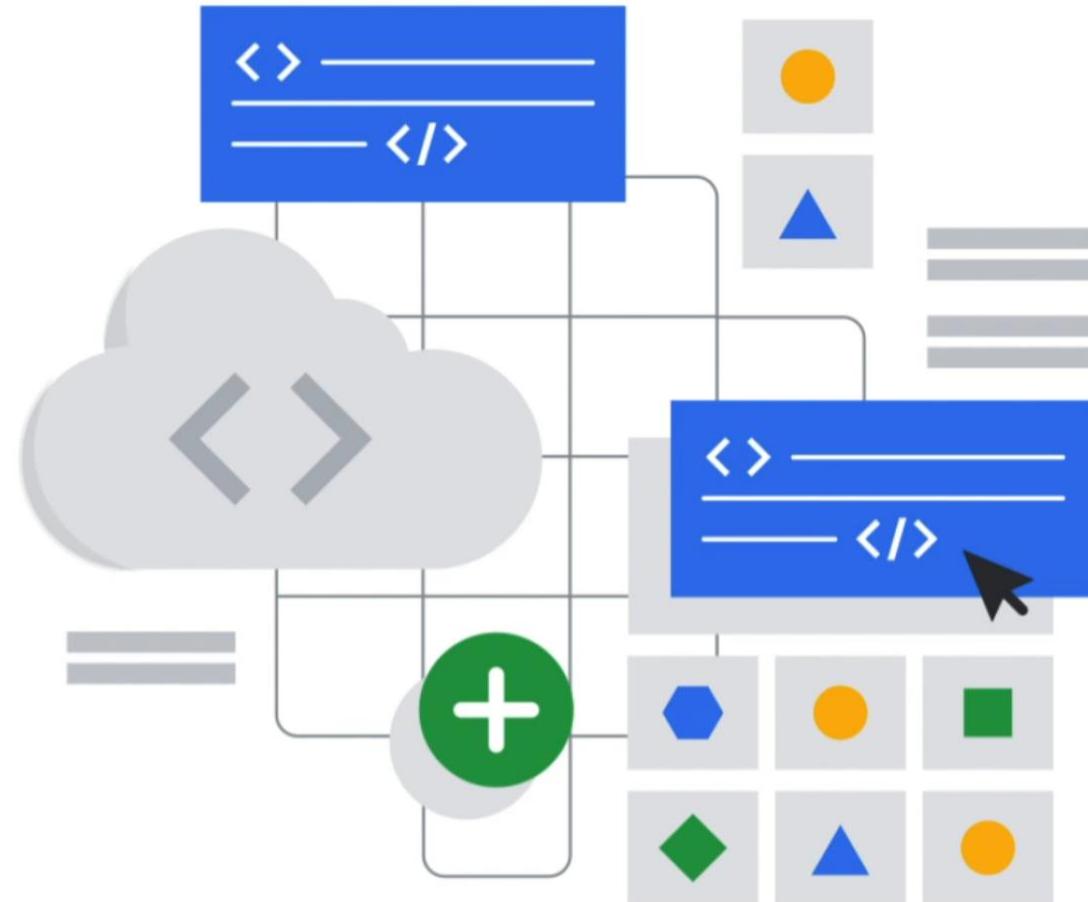


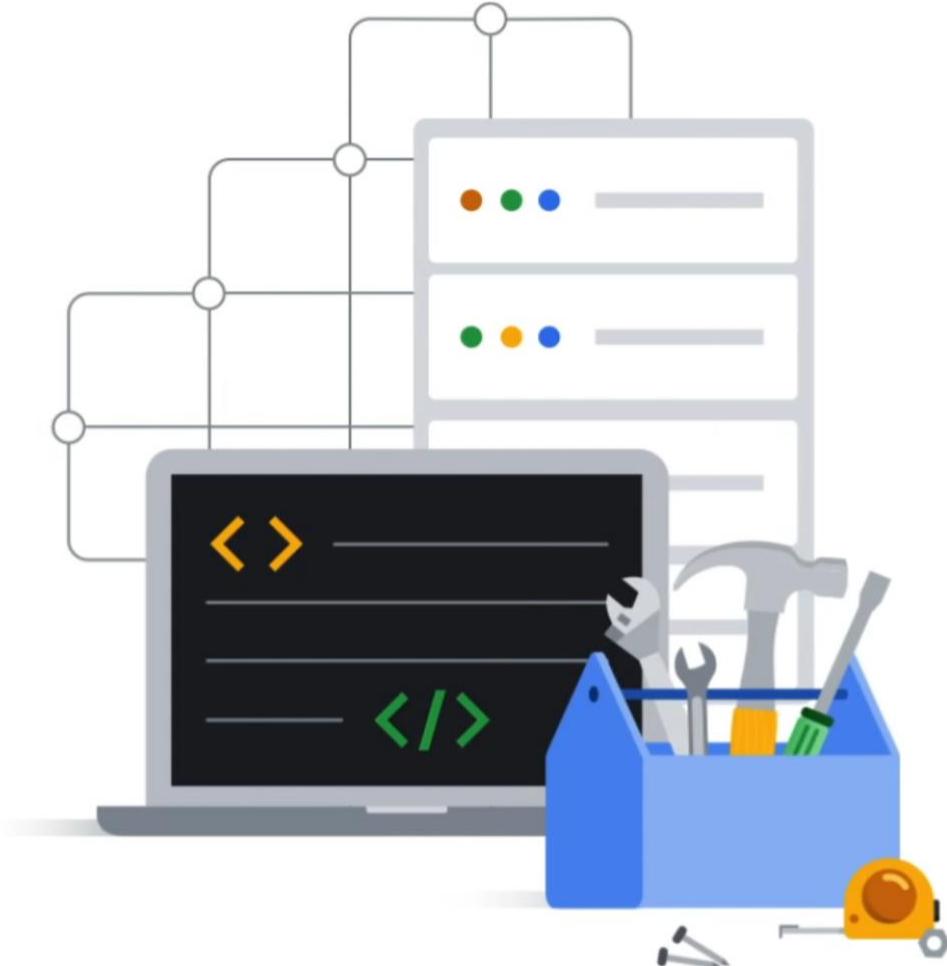
# Introduction to Terraform for Google Cloud



# What is infrastructure as code (IaC)?

Instead of clicking around a web UI or using SSH to connect to a server and manually executing commands, **with IaC you can write code in files to define, provision, and manage your infrastructure.**





## DevOps

Emphasizes the collaboration and communication of both software developers and IT operations teams

Automate the process of software delivery and infrastructure changes

# Problems that IaC can solve



Inability to scale rapidly

Requires rapid scaling of IT infrastructure



Operational bottlenecks

Challenge of managing infrastructure consistently in scale



Disconnected feedback loops

Communication gap between software and IT teams



High manual errors

Increased scale leads to increased human errors

# Benefits of IaC

Declarative	Specify the desired state of infrastructure, not updates.
Code management	Commit, version, trace, and collaborate, just like source code.
Auditable	Compare infrastructure between desired state and current state.
Portable	Build reusable modules across an organization.

# Provisioning versus configuration

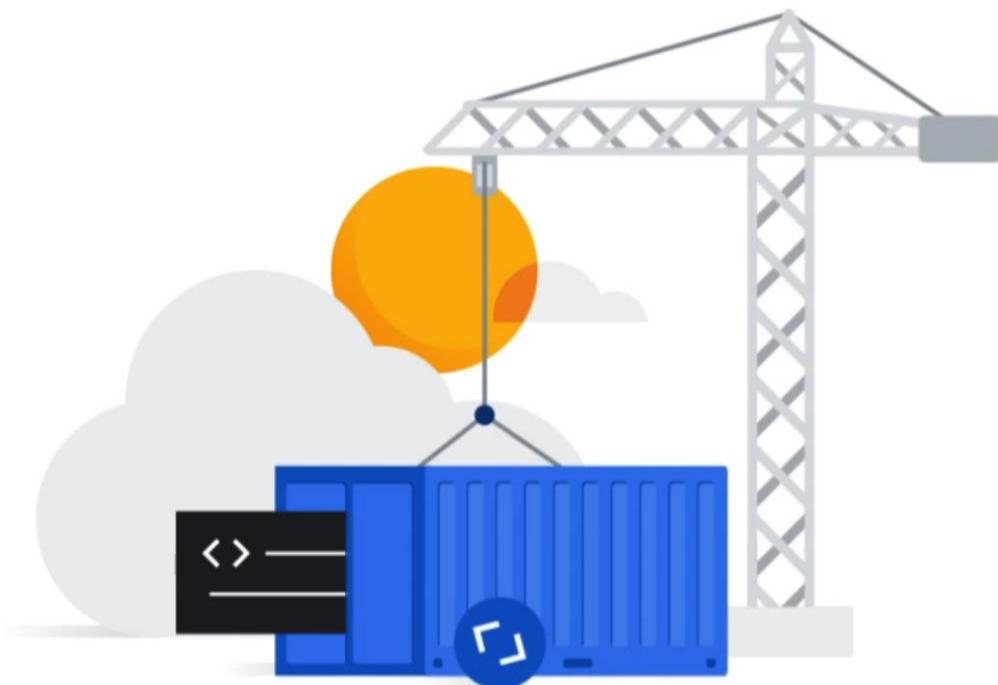
## Infrastructure as code

- Used for provisioning and managing cloud resources.
- Example: Creating and provisioning a VM instance.
- Referring to frameworks that manipulate Google Cloud APIs to deploy the infrastructure.

## Configuration Management

- Used for virtual machine OS-level configuration.
- Example: Configuring the internals of the VMs.
- Referring to package configurations and software maintenance.

# Provisioning versus configuration



Infrastructure as a code

- Launch a GKE cluster

Configuration management

- Deploy containers into the GKE cluster

# IaC takes a declarative approach to infrastructure

Imperative (command)

Command line

“Give me five servers”

How to create?

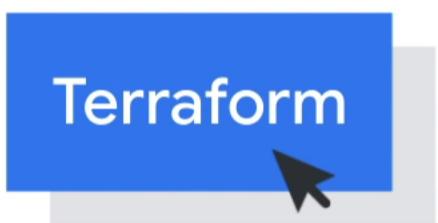
Declarative (statement)

YAML

“I should have five servers”

What to create?

VS.



**Terraform** is an open source infrastructure as code tool created by HashiCorp that lets you provision Google Cloud resources with **declarative** configuration files

## Terraform features



Multi-cloud and multi-API



Open core with enterprise support



Large community

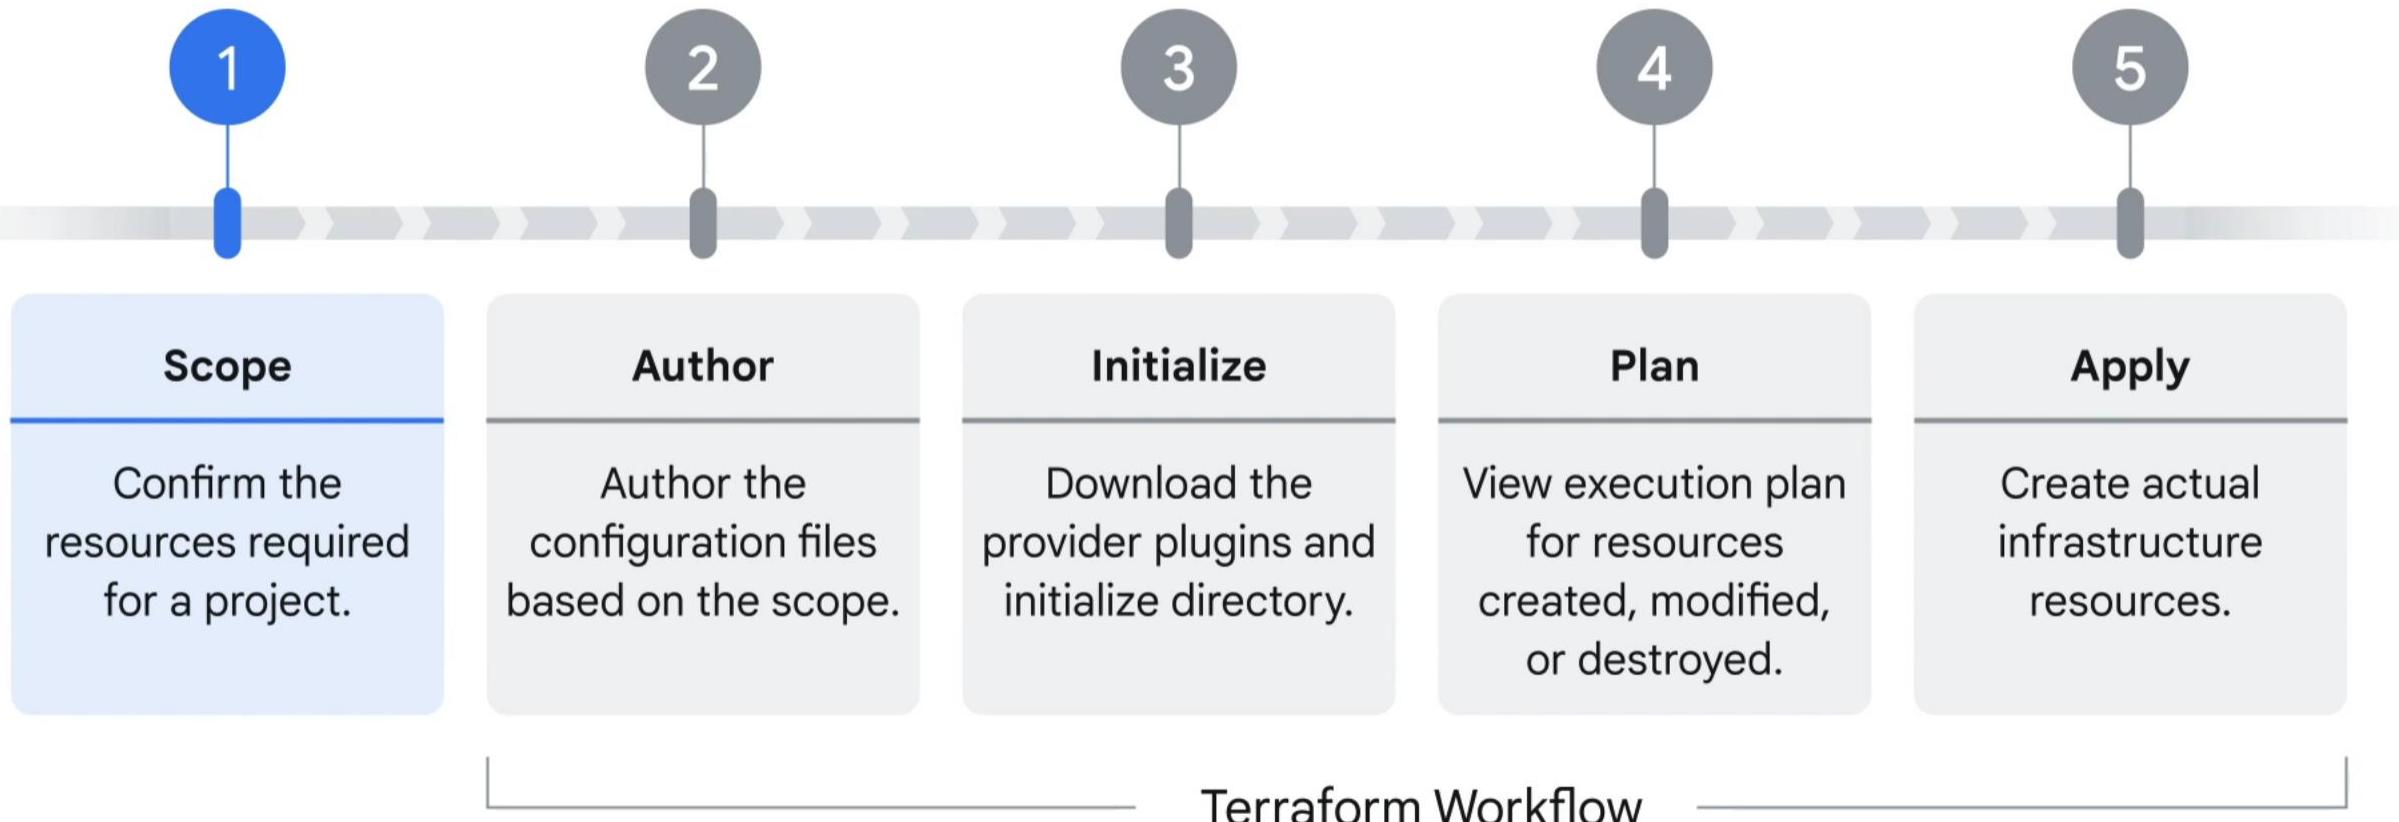


Infrastructure provisioning

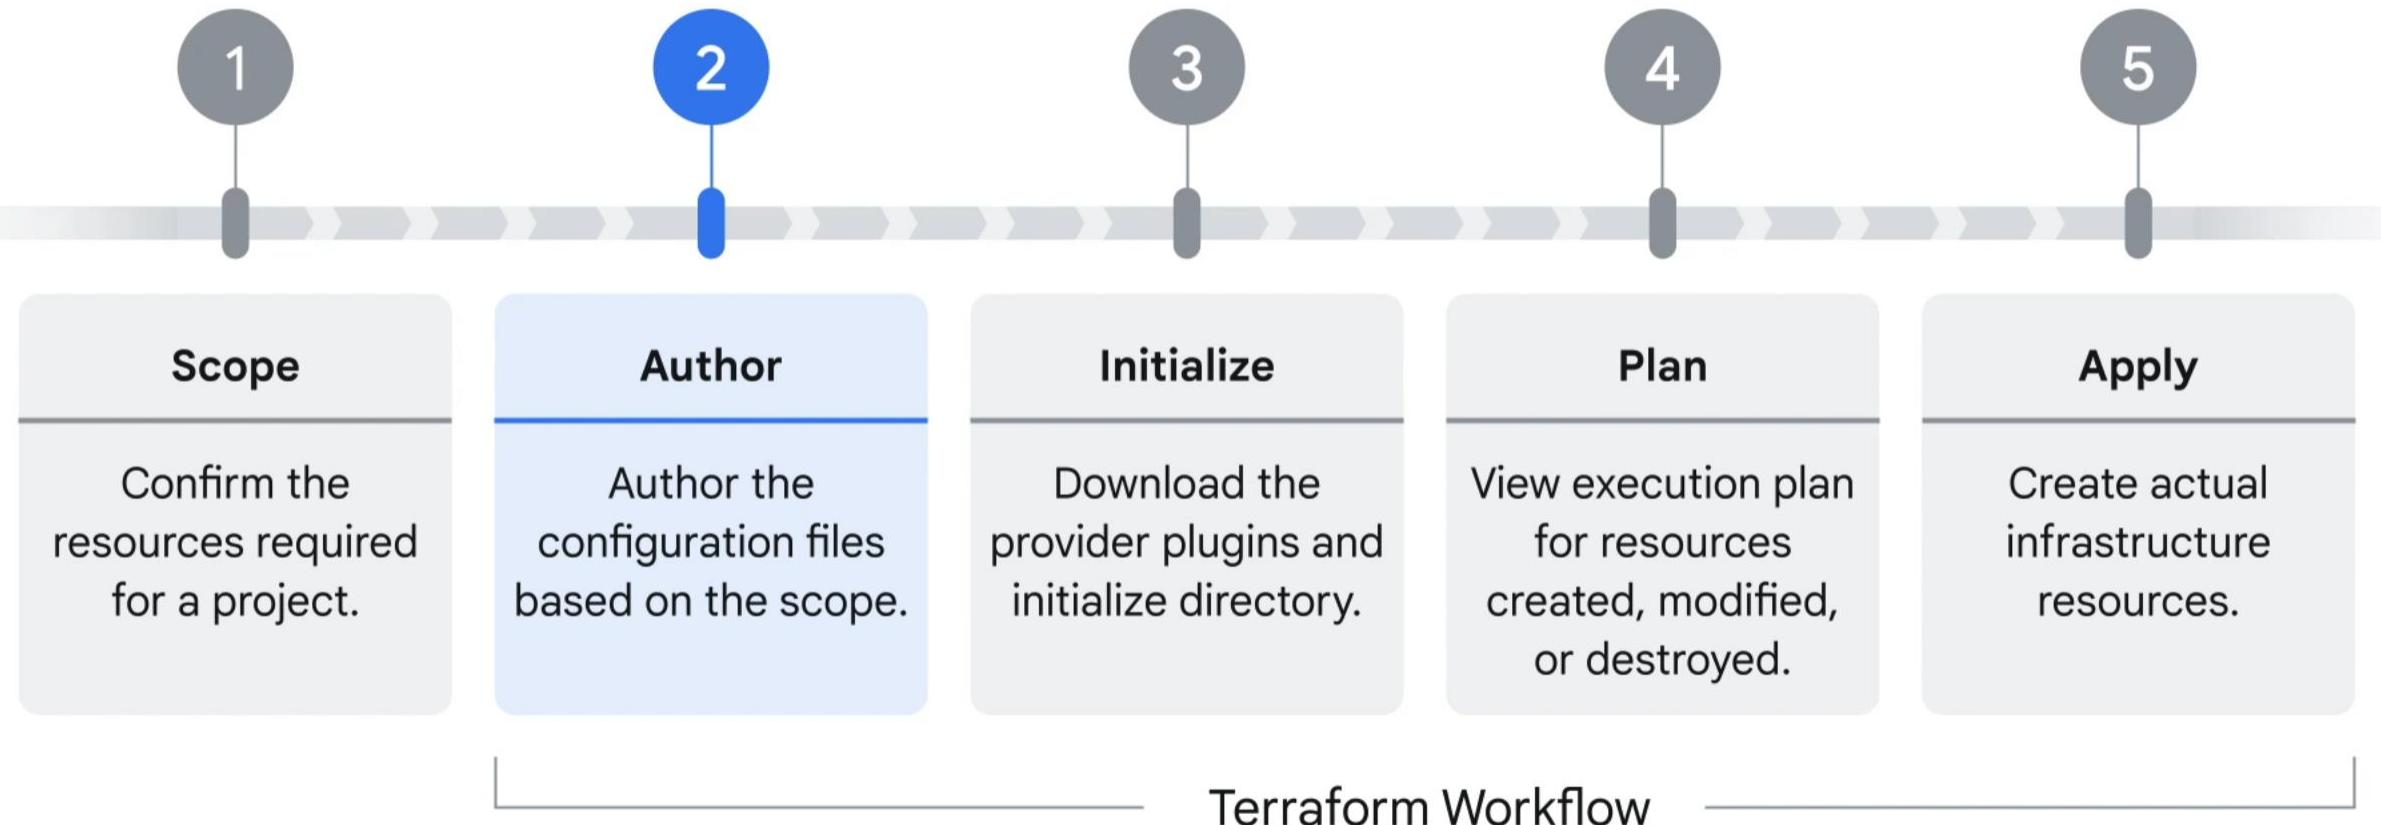
# Terraform for Google Cloud

-  Provision resources
-  Create resource dependencies
-  Standardize configurations
-  Validate inputs to resource arguments

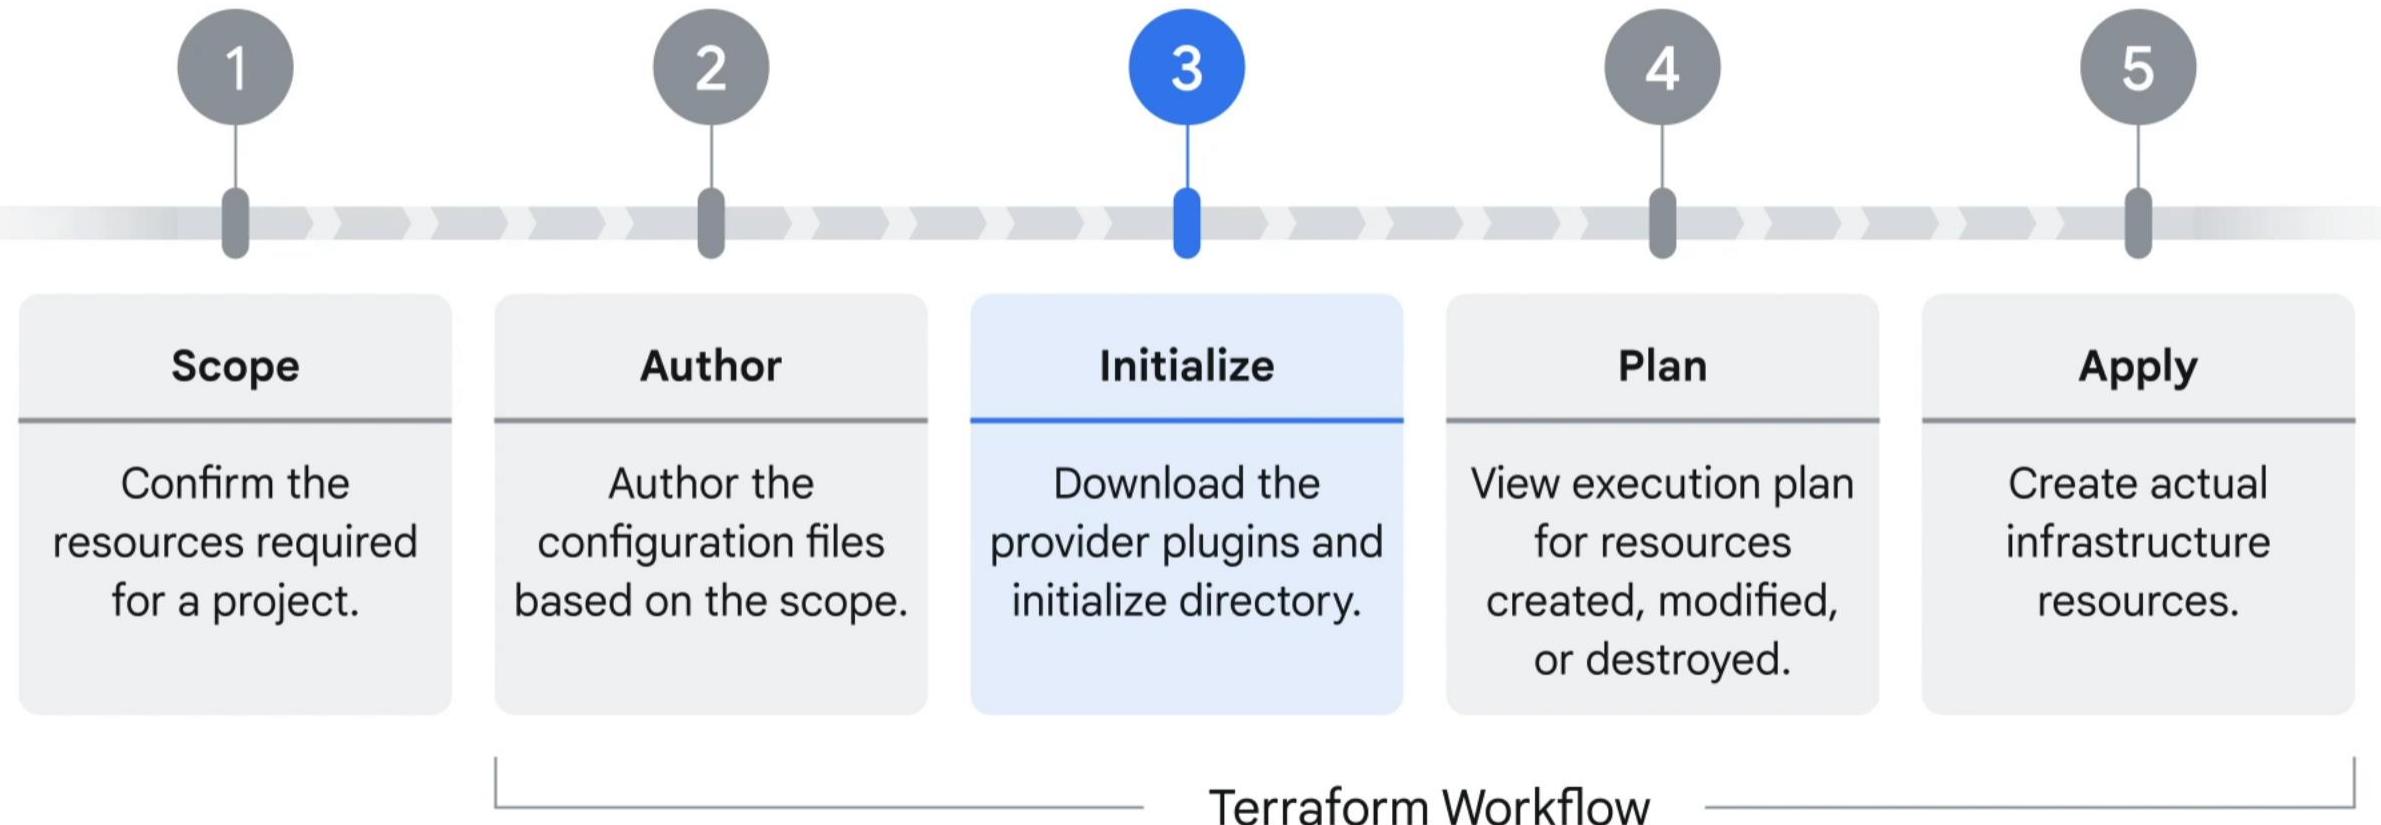
# IaC configuration workflow



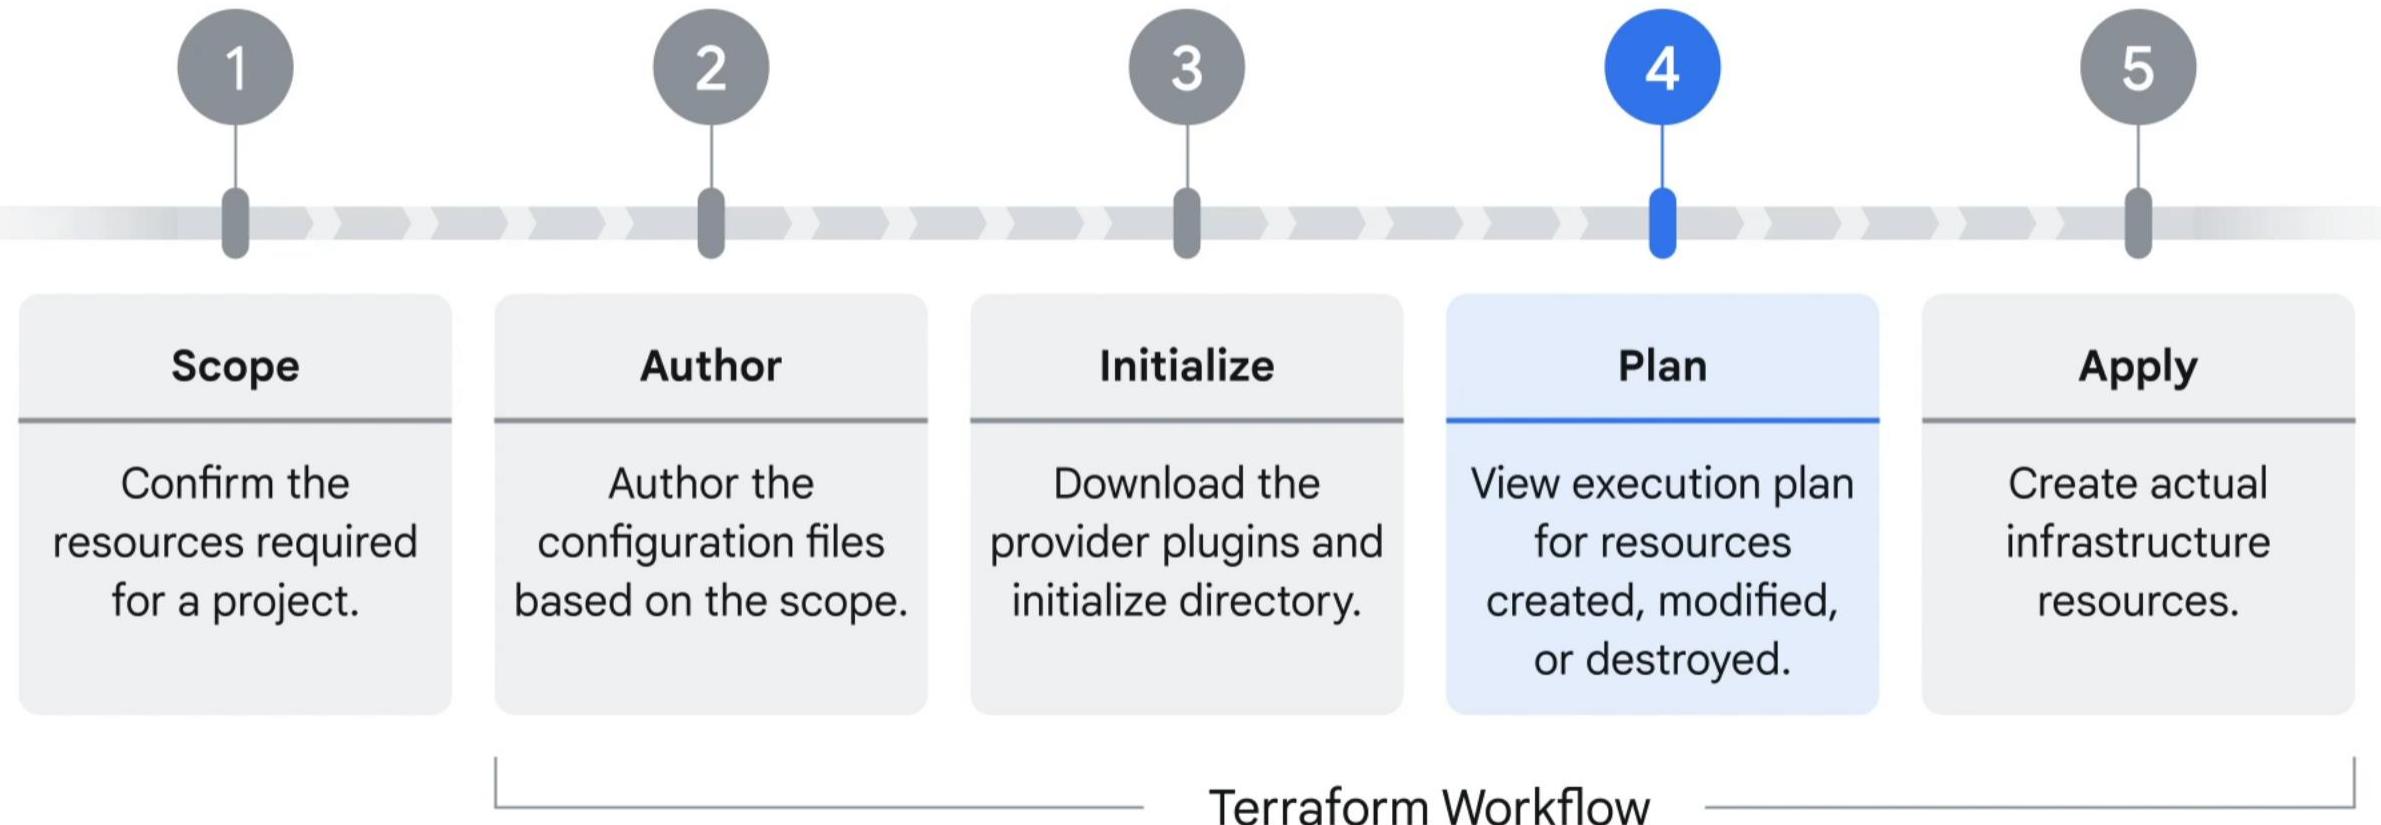
# IaC configuration workflow



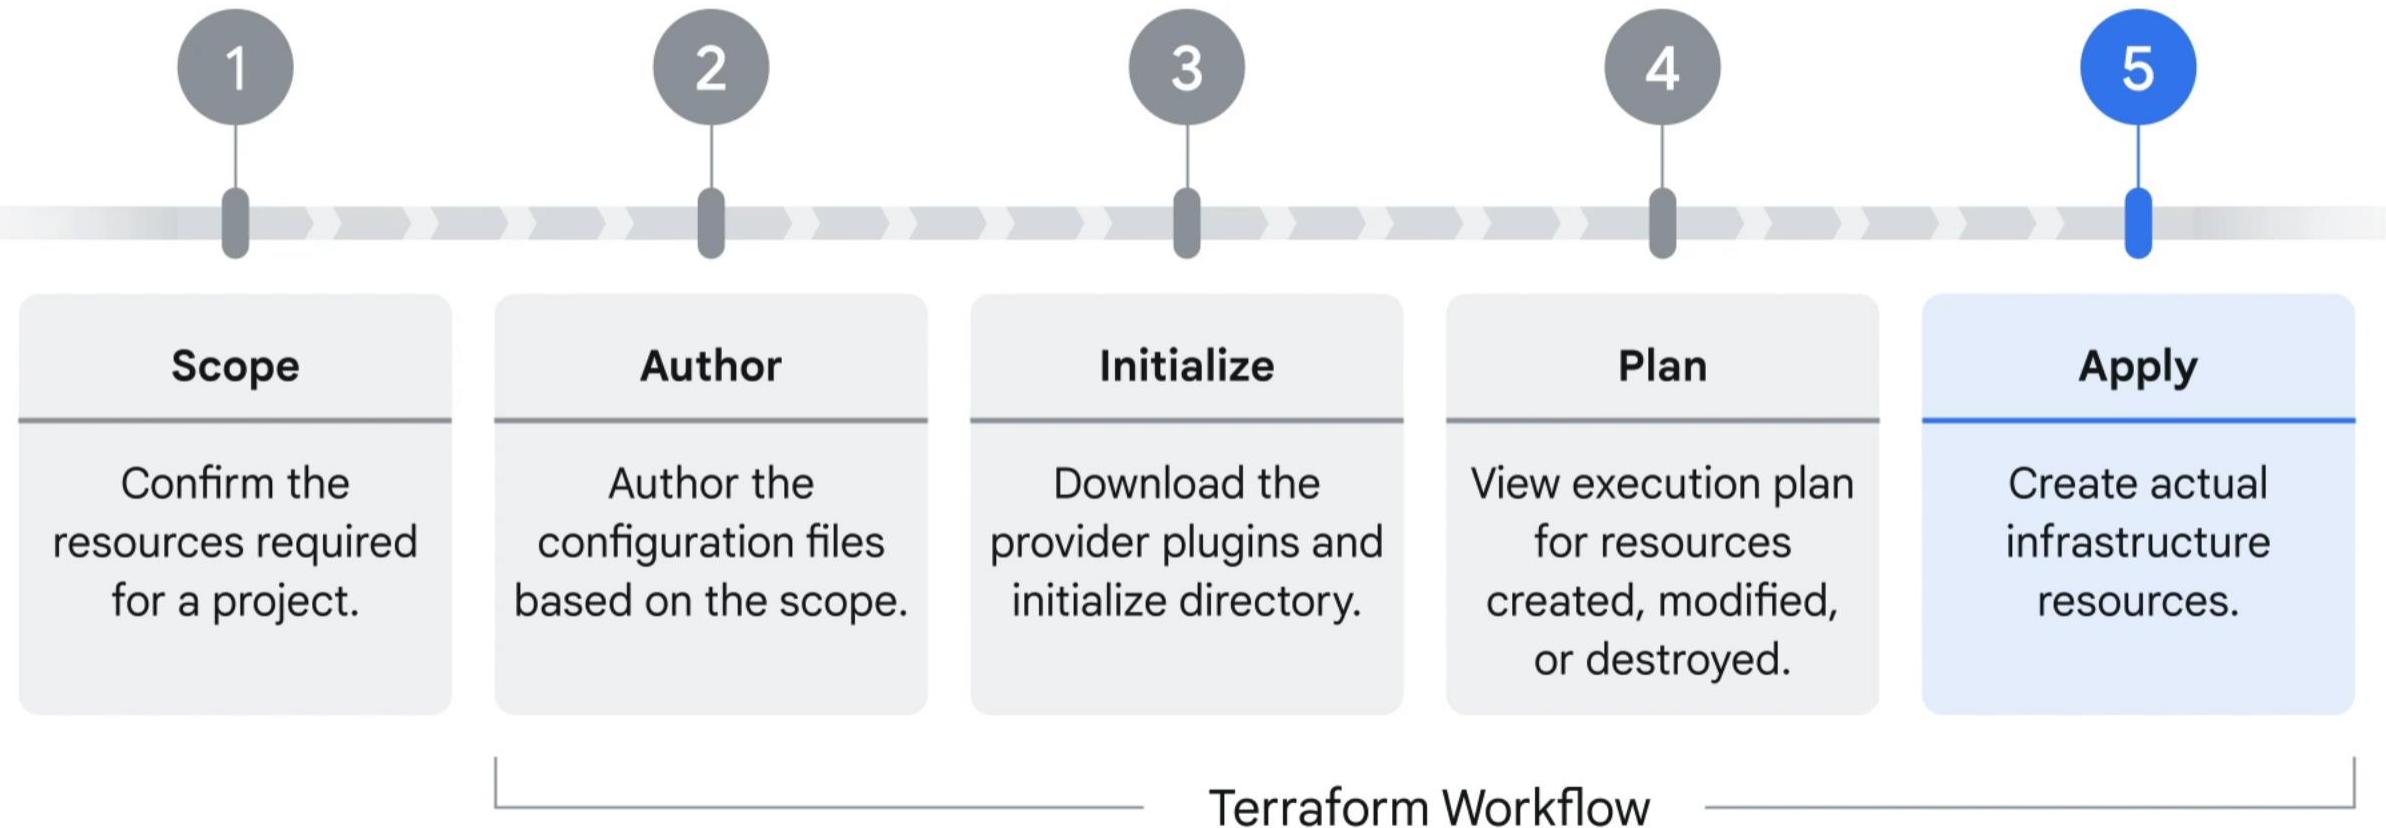
# IaC configuration workflow



# IaC configuration workflow



# IaC configuration workflow

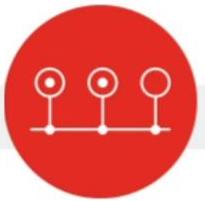


# Terraform use cases



## Manage infrastructure

Terraform takes an immutable approach to building and managing infrastructure.



## Track changes

Terraform enables you to review the changes before they are applied to the configuration setup.



## Automate changes

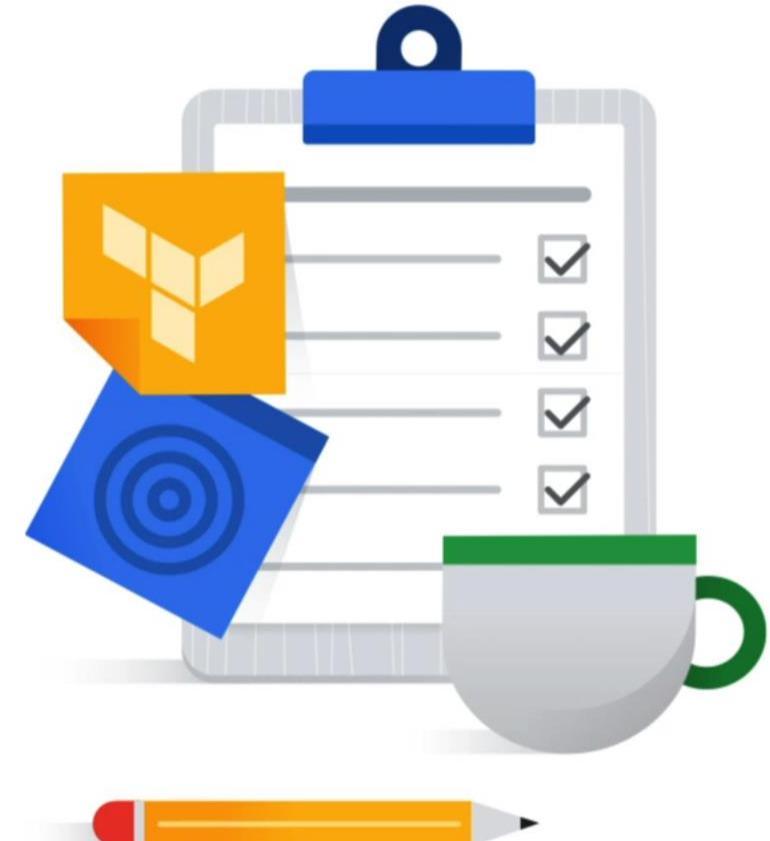
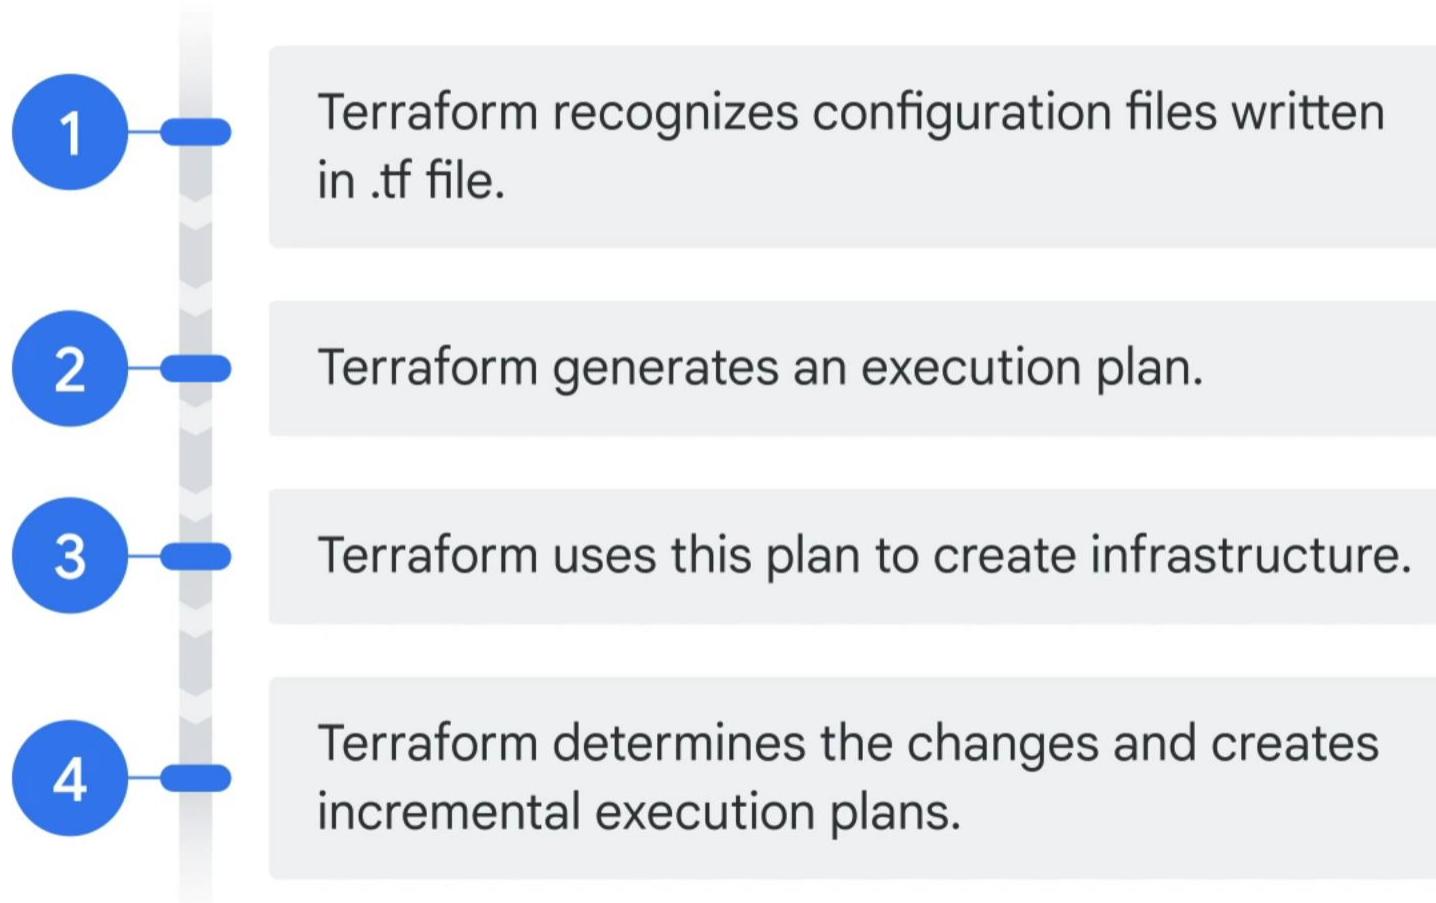
Terraform defines the end state of the infrastructure instead of a series of steps to achieve it.



## Standardize the configuration

Terraform uses modules to implement best practices and improve efficiency.

# Using Terraform



# Running Terraform in production

	Managed	Pros	Cons
Terraform Open Source		<ul style="list-style-type: none"><li>• Deployed on a local machine or compute resource in the cloud</li><li>• No license cost</li><li>• Use public registry within your code</li></ul>	<ul style="list-style-type: none"><li>• Does not support concurrent deployments</li><li>• Only interfaced through CLI</li></ul>
Terraform Cloud		<ul style="list-style-type: none"><li>• SaaS based version</li><li>• Small operational overhead</li><li>• Comes with three plans</li><li>• Supports concurrent deployments</li><li>• Can be accessed through GUI and CLI</li></ul>	<ul style="list-style-type: none"><li>• License cost for advanced features</li></ul>
Terraform Enterprise		<ul style="list-style-type: none"><li>• Private implementation</li><li>• Supports concurrent deployments</li><li>• Secure deployment</li><li>• Can be accessed through GUI and CLI</li></ul>	<ul style="list-style-type: none"><li>• Infrastructure and license costs</li><li>• Large operational overhead</li></ul>

# Running Terraform in production

	Managed	Pros	Cons
Terraform Open Source		<ul style="list-style-type: none"><li>• Deployed on a local machine or compute resource in the cloud</li><li>• No license cost</li><li>• Use public registry within your code</li></ul>	<ul style="list-style-type: none"><li>• Does not support concurrent deployments</li><li>• Only interfaced through CLI</li></ul>
Terraform Cloud		<ul style="list-style-type: none"><li>• SaaS based version</li><li>• Small operational overhead</li><li>• Comes with three plans</li><li>• Supports concurrent deployments</li><li>• Can be accessed through GUI and CLI</li></ul>	<ul style="list-style-type: none"><li>• License cost for advanced features</li></ul>
Terraform Enterprise		<ul style="list-style-type: none"><li>• Private implementation</li><li>• Supports concurrent deployments</li><li>• Secure deployment</li><li>• Can be accessed through GUI and CLI</li></ul>	<ul style="list-style-type: none"><li>• Infrastructure and license costs</li><li>• Large operational overhead</li></ul>

# Running Terraform in production

	Managed	Pros	Cons
Terraform Open Source		<ul style="list-style-type: none"><li>• Deployed on a local machine or compute resource in the cloud</li><li>• No license cost</li><li>• Use public registry within your code</li></ul>	<ul style="list-style-type: none"><li>• Does not support concurrent deployments</li><li>• Only interfaced through CLI</li></ul>
Terraform Cloud		<ul style="list-style-type: none"><li>• SaaS based version</li><li>• Small operational overhead</li><li>• Comes with three plans</li><li>• Supports concurrent deployments</li><li>• Can be accessed through GUI and CLI</li></ul>	<ul style="list-style-type: none"><li>• License cost for advanced features</li></ul>
Terraform Enterprise		<ul style="list-style-type: none"><li>• Private implementation</li><li>• Supports concurrent deployments</li><li>• Secure deployment</li><li>• Can be accessed through GUI and CLI</li></ul>	<ul style="list-style-type: none"><li>• Infrastructure and license costs</li><li>• Large operational overhead</li></ul>

# Installing Terraform on local machine

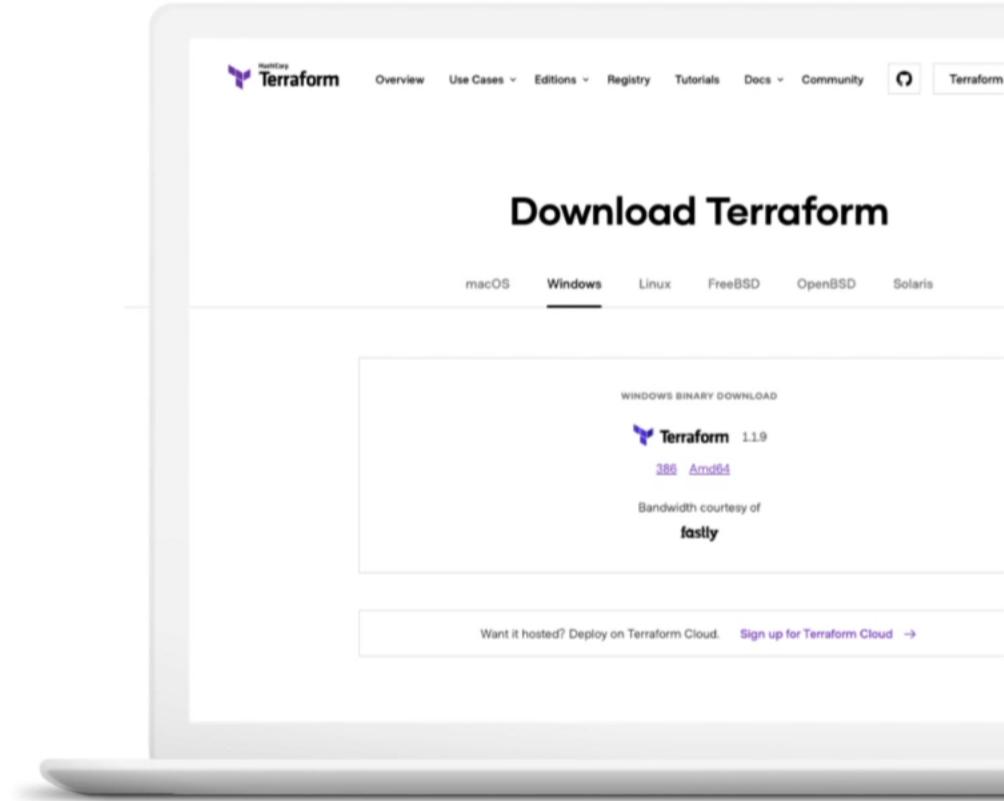
- 1
- 2
- 3
- 4

Download the appropriate package based on your system.

Unzip the package. Terraform includes a single binary called **terraform**.

Edit the **PATH** variable to include **terraform**.

To verify the installation, enter **terraform -help** in a new terminal.



# Authentication for Google Cloud



## On your workstation

Authenticate Terraform using Google Cloud SDK.  
On Cloud Shell, it is pre-authenticated for you.



## In a VM on Google Cloud

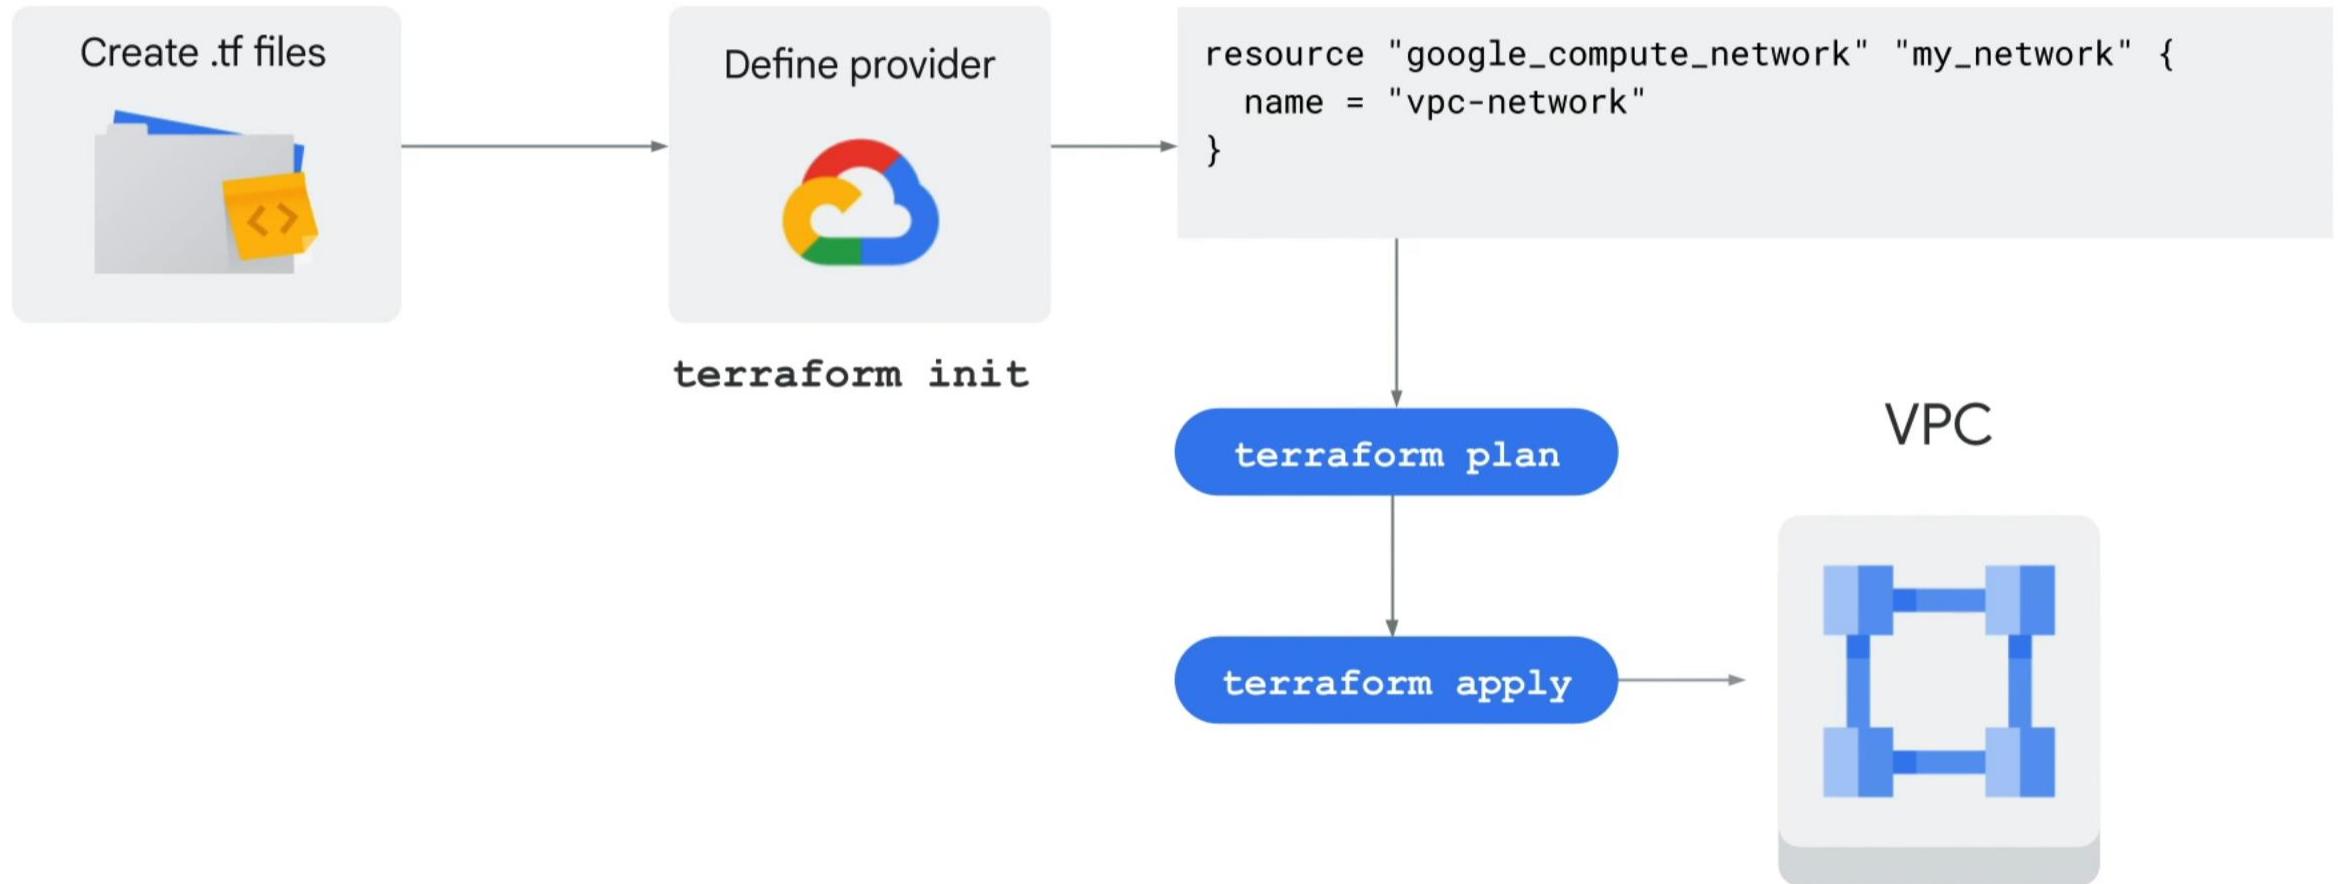
Configure the VM to use a Google Service Account.



## Outside Google Cloud

Using workload identity federation, generate a service account key and set environment variables.

# Example: Creating a VPC network



# Terraform directory

- Terraform uses configuration files to declare an infrastructure element.
- The configuration is written in terraform language with a .tf extension.
- A configuration consists of:
  - A root module/ root configuration
  - Zero or more child modules
  - Variable.tf (optional but recommended)
  - Outputs.tf (optional but recommended)
- Terraform commands are run on the working directory.

```
-- main.tf  
-- servers/  
    -- main.tf  
    -- providers.tf  
    -- variables.tf  
    -- outputs.tf
```



Root module

Child module

# HashiCorp Configuration Language (HCL)

## Syntax

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {
    # Block body
    <IDENTIFIER> = <EXPRESSION> #Argument
}
```

# HashiCorp Configuration Language (HCL)

- Terraform's configuration language for creating and managing API-based resources
- Configuration language, not a programming language.
- Includes limited set of primitives such as variables, resources, outputs and modules.
- Includes no traditional statements or control loops.

## Syntax

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
    # Block body  
    <IDENTIFIER> = <EXPRESSION> #Argument  
}
```

# HCL syntax

Blocks

Arguments

Identifiers

Expressions

Comments

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
    # Block body  
    <IDENTIFIER> = <EXPRESSION> #Argument  
}
```

```
resource "google_compute_network" "default" {  
    # custom mode network definition  
    name                  = mynetwork  
    auto_create_subnetworks = false  
}
```

# HCL syntax

Blocks

Arguments

Identifiers

Expressions

Comments

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
    # Block body  
    <IDENTIFIER> = <EXPRESSION> #Argument  
}
```

```
resource "google_compute_network" "default" {  
    # custom mode network definition  
    name                  = mynetwork  
    auto_create_subnetworks = false  
}
```

# HCL syntax

Blocks

Arguments

Identifiers

Expressions

Comments

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
    # Block body  
    <IDENTIFIER> = <EXPRESSION> #Argument  
}
```

```
resource "google_compute_network" "default" {  
    # custom mode network definition  
    name                  = mynetwork  
    auto_create_subnetworks = false  
}
```

# HCL syntax

Blocks

Arguments

Identifiers

Expressions

Comments

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
    # Block body  
    <IDENTIFIER> = <VALUE/EXPRESSION> #Argument  
}
```

```
resource "google_compute_network" "default" {  
    # custom mode network definition  
    name                  = mynetwork  
    auto_create_subnetworks = false  
}
```

# HCL syntax

Blocks

Arguments

Identifiers

Expressions

Comments

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
    # Block body  
    <IDENTIFIER> = <EXPRESSION> #Argument  
}
```

```
resource "google_compute_network" "default" {  
    # custom mode network definition  
    name                  = mynetwork  
    auto_create_subnetworks = false  
}
```

# Resources

```
-- main.tf  
-- providers.tf  
-- variables.tf  
-- outputs.tf
```

- Resources are code blocks that define the infrastructure components.

Example: Cloud Storage Bucket

- Terraform uses the **resource type** and the **resource name** to identify an infrastructure element

```
resource "resource_type" "resource_name" {  
    #Resource specific arguments  
}
```

# Examples of resources

```
-- main.tf
-- providers.tf
-- variables.tf
-- outputs.tf
```

- The keyword **resource** is used to identify the block as the cloud infrastructure component.

```
resource "google_storage_bucket" "example-bucket" {
  name        = "<unique-bucket-name>"
  location    = "US"
}
```

# Examples of resources

```
-- main.tf
-- providers.tf
-- variables.tf
-- outputs.tf
```

- The keyword **resource** is used to identify the block as the cloud infrastructure component.
- Resource type is **google\_storage\_bucket**.
- The resource name is **example-bucket**.

```
resource "google_storage_bucket" "example-bucket" {
  name        = "<unique-bucket-name>"
  location    = "US"
}
```

# Resource arguments

```
-- main.tf
-- providers.tf
-- variables.tf
-- outputs.tf
```

- The arguments differ based on the resource type.
- Some arguments are required, others are optional.

```
resource "google_storage_bucket" "example-bucket" {
  name        = "<unique-bucket-name>" //Required
  location    = "US"
}

resource "google_compute_instance" "my_instance" {
  name        = "test"
  machine_type = "e2-medium"
  zone        = "us-central1-a"
  boot_disk {
    initialize_params {
      image = "debian-cloud/debian-9"
    }
  }
}
```

# Providers

```
-- main.tf
-- providers.tf
-- variables.tf
-- outputs.tf
```

- Terraform downloads the provider plugin in the root configuration when the provider is declared.
- Providers expose specific APIs as Terraform resources and manage their interactions.

```
terraform {
  required_providers {
    google = {
      source = "hashicorp/google"
      version = "4.23.0"
    }
  }
}

provider "google" {
  # Configuration options
  project = <project_id>
  region  = "us-central1"
}
```

# Providers

```
-- main.tf  
-- providers.tf  
-- variables.tf  
-- outputs.tf
```

- Provider configurations belong in the root module of a Terraform configuration.
- Arguments such as project and region can be declared within the provider block.

```
terraform {  
  required_providers {  
    google = {  
      source  = "hashicorp/google"  
      version = "4.23.0"  
    }  
  }  
  
  provider "google" {  
    # Configuration options  
    project = <project_id>  
    region  = "us-central1"  
  }  
}
```

# Provider versions

```
-- main.tf
-- providers.tf
-- variables.tf
-- outputs.tf
```

- The version argument is optional, but recommended.
- The version argument is used to constrain the provider to a specific version or a range of versions.

```
terraform {
  required_providers {
    google = {
      source = "hashicorp/google"
      version = "4.23.0"
    }
  }
}

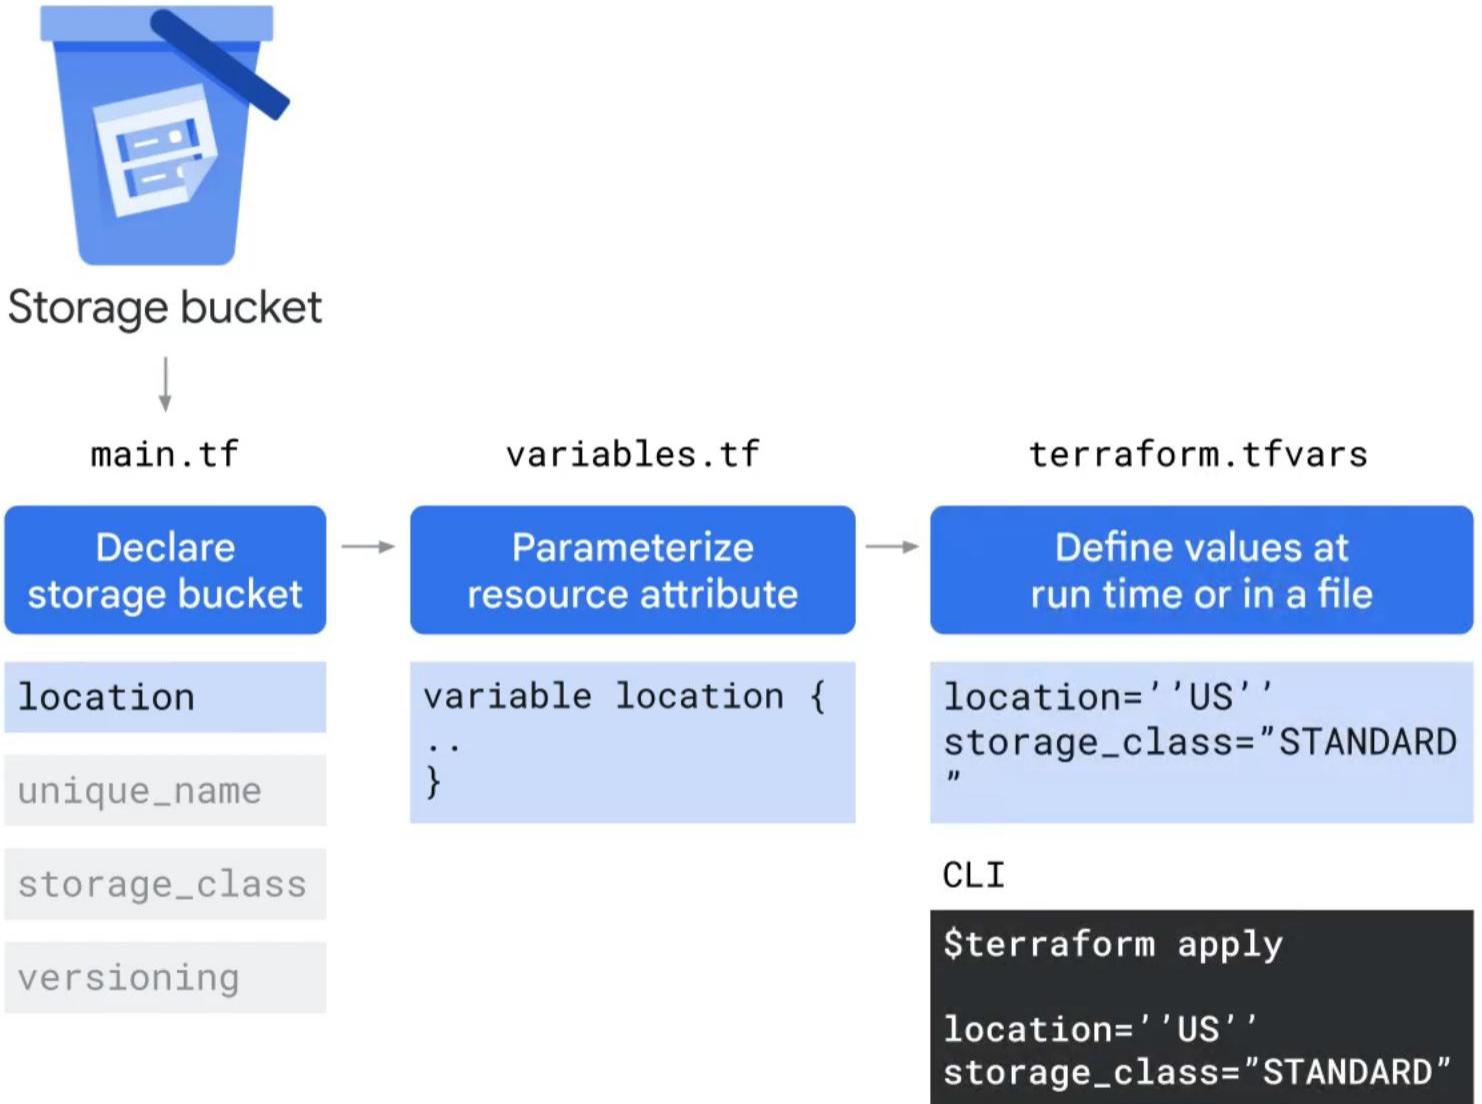
provider "google" {
  # Configuration options
}
```

# Variables

- Parameterize resource arguments to eliminate hard coding its values

Example: Region, project ID, zone, etc.

- Define a resource attribute at run time or centrally in a file with a .tfvars extension.



# Outputs

```
-- main.tf  
-- providers.tf  
-- variables.tf  
-- outputs.tf
```

- Output values are stored in outputs.tf file.
- Output values expose values of resource attributes.

```
output "bucket_URL" {  
    value = google_storage_bucket.mybucket.URL  
}
```

```
#terraform apply  
Google_storage_bucket.mybucket: Creating...  
Google_storage_bucket.mybucket: Creating complete after 1s []  
Apply complete! Resources: 1 added, 0 changed, 0 destroyed.  
Outputs:  
bucket_URL = "https://storage.googleapis.com/my-gallery/..
```

# State

```
-- main.tf  
-- providers.tf  
-- variables.tf  
-- outputs.tf  
-- terraform.tfstate
```

- Terraform saves the state of resources it manages in a state file.
- The state file can be stored:
  - Locally (default)
  - Remotely in a shared location

**You do not modify this file.**

```
{  
  "version": 4,  
  "terraform_version": "1.0.11",  
  "serial": 3,  
  "lineage": "822c3d96-0500-29cd-68e3-13101f2846f0",  
  "outputs": {  
    "vm_name": {  
      "value": "terraform-test",  
      "type": "string"  
    }  
  },  
  "resources": [  
    {  
      "mode": "managed",  
      "type": "google_compute_instance",  
      "name": "default",  
      "provider":  
        "provider": \"registry.terraform.io/hashicorp/google\"",  
      "instances": [  
        {
```

# Modules

A Terraform module is a set of Terraform configuration files in a single directory.

- It is the primary method for code reuse in Terraform.
- There are 2 kinds of sources:
  - Local: Source within your directory
  - Remote: Source outside your directory.

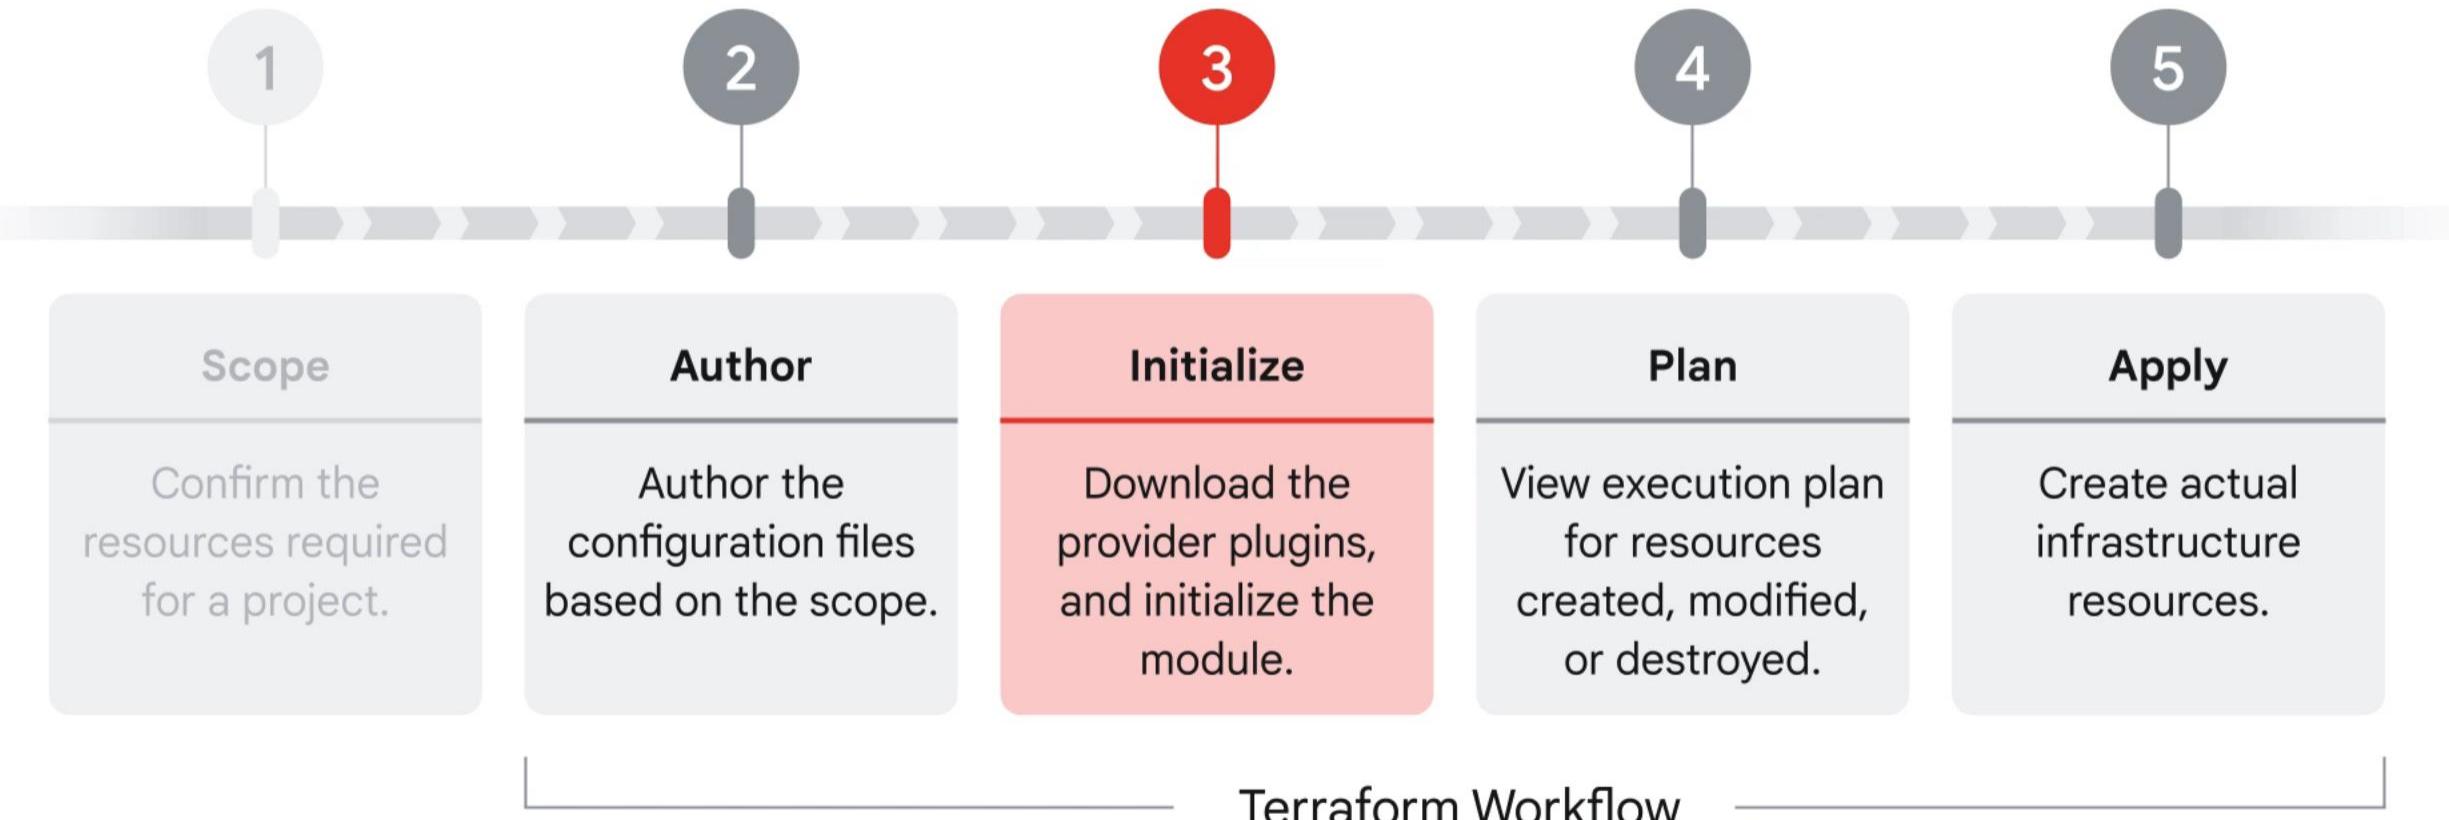
```
-- instances/  
-- main.tf  
-- variables.tf  
-- outputs.tf
```

modules

# Terraform commands

<code>terraform init</code>	Initialize the provider with plugin
<code>terraform plan</code>	Preview of resources that will be created after terraform apply
<code>terraform apply</code>	Create real infrastructure resources
<code>terraform destroy</code>	Destroy infrastructure resources
<code>terraform fmt</code>	Auto format to match canonical conventions

# Initialize Terraform using `terraform init`



# Initialize Phase:

`terraform init` downloads the provider plugins

```
-- main.tf
-- servers/
  -- main.tf
  -- variables.tf
  -- outputs.tf
}
}

  terraform {
    required_providers {
      google = {
        source = "hashicorp/google"
      }
    }
  }
```

# Initialize Phase:

`terraform init` downloads the provider plugins

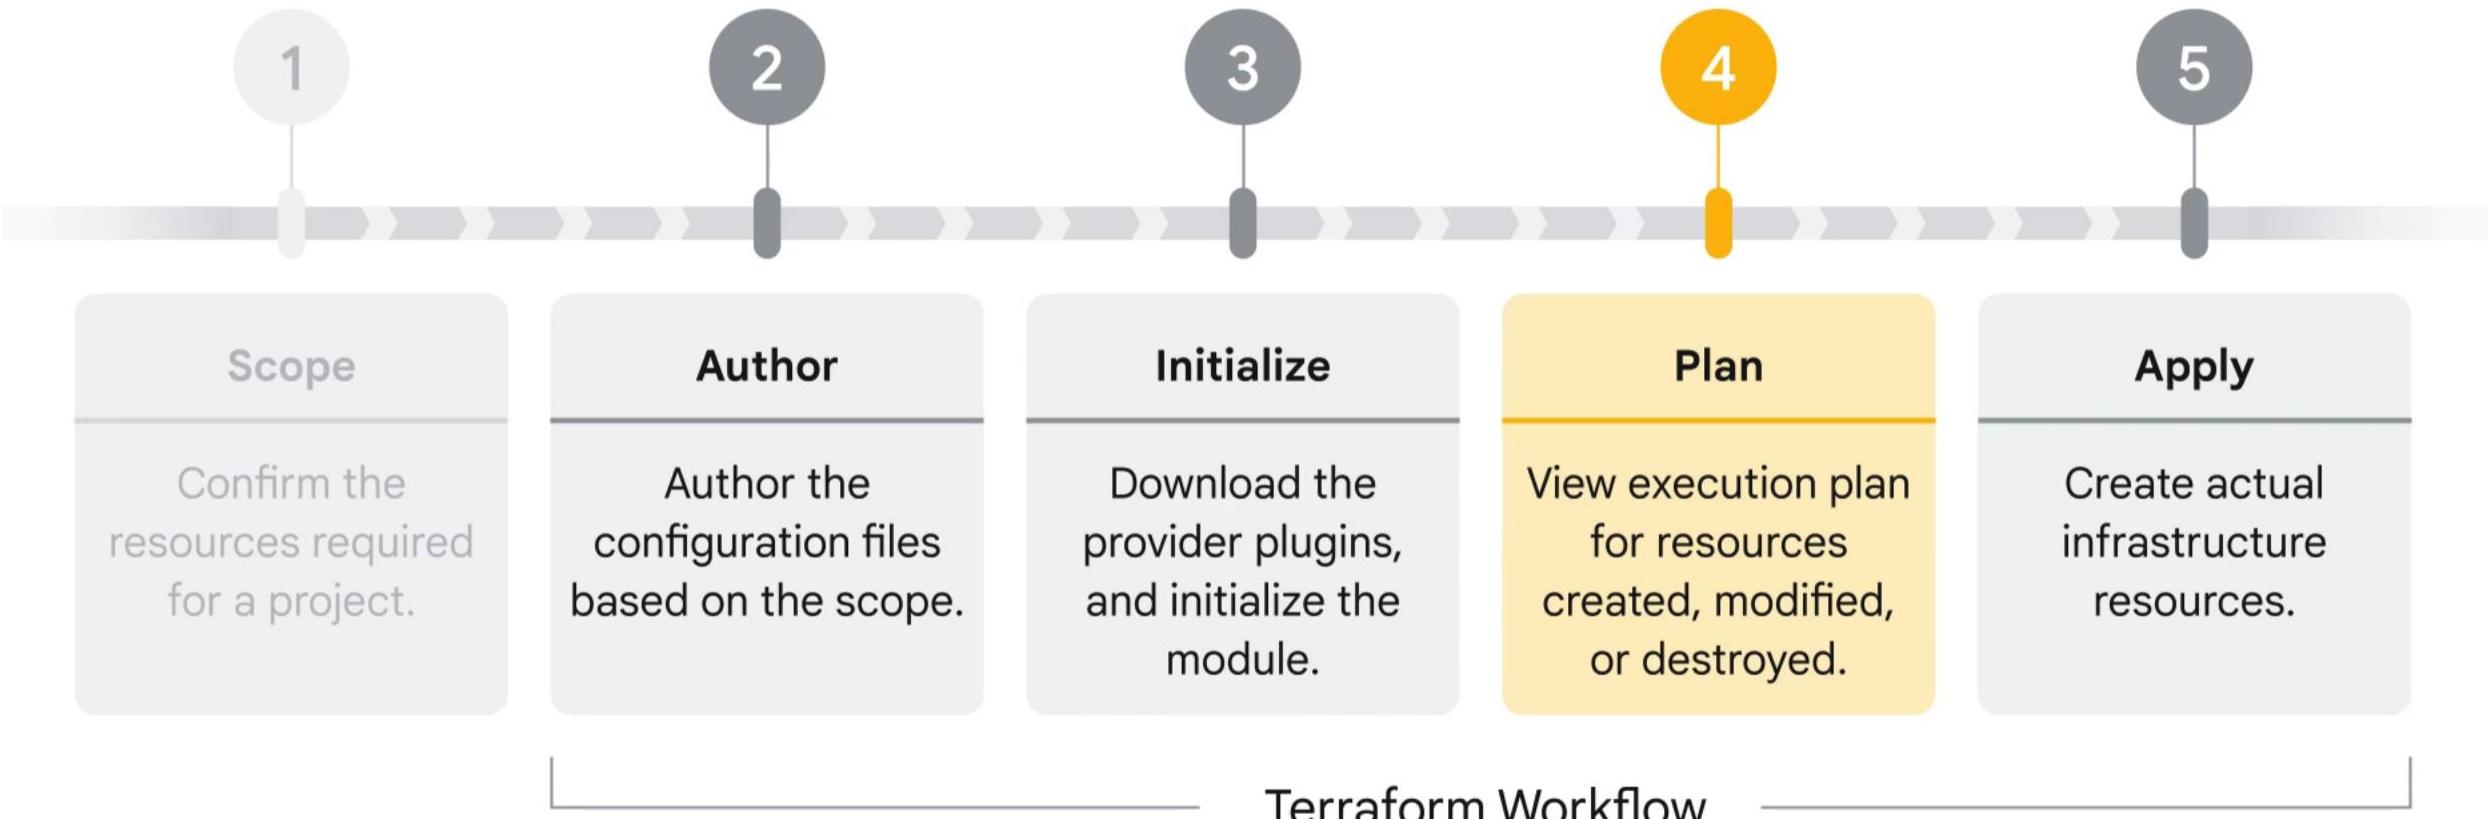
```
-- main.tf
-- servers/
  -- main.tf
  -- variables.tf
  -- outputs.tf
```

```
  terraform {
    required_providers {
      google = {
        source = "hashicorp/google"
      }
    }
  }
```

```
$ terraform init
Initializing the backend...
```

```
Initializing provider plugins...
- Finding latest version of hashicorp/google...
- Installing hashicorp/google v4.21.0...
- Installed hashicorp/google v4.21.0 (signed by HashiCorp)
..
Terraform has been successfully initialized!
```

# Preview resource action using terraform plan



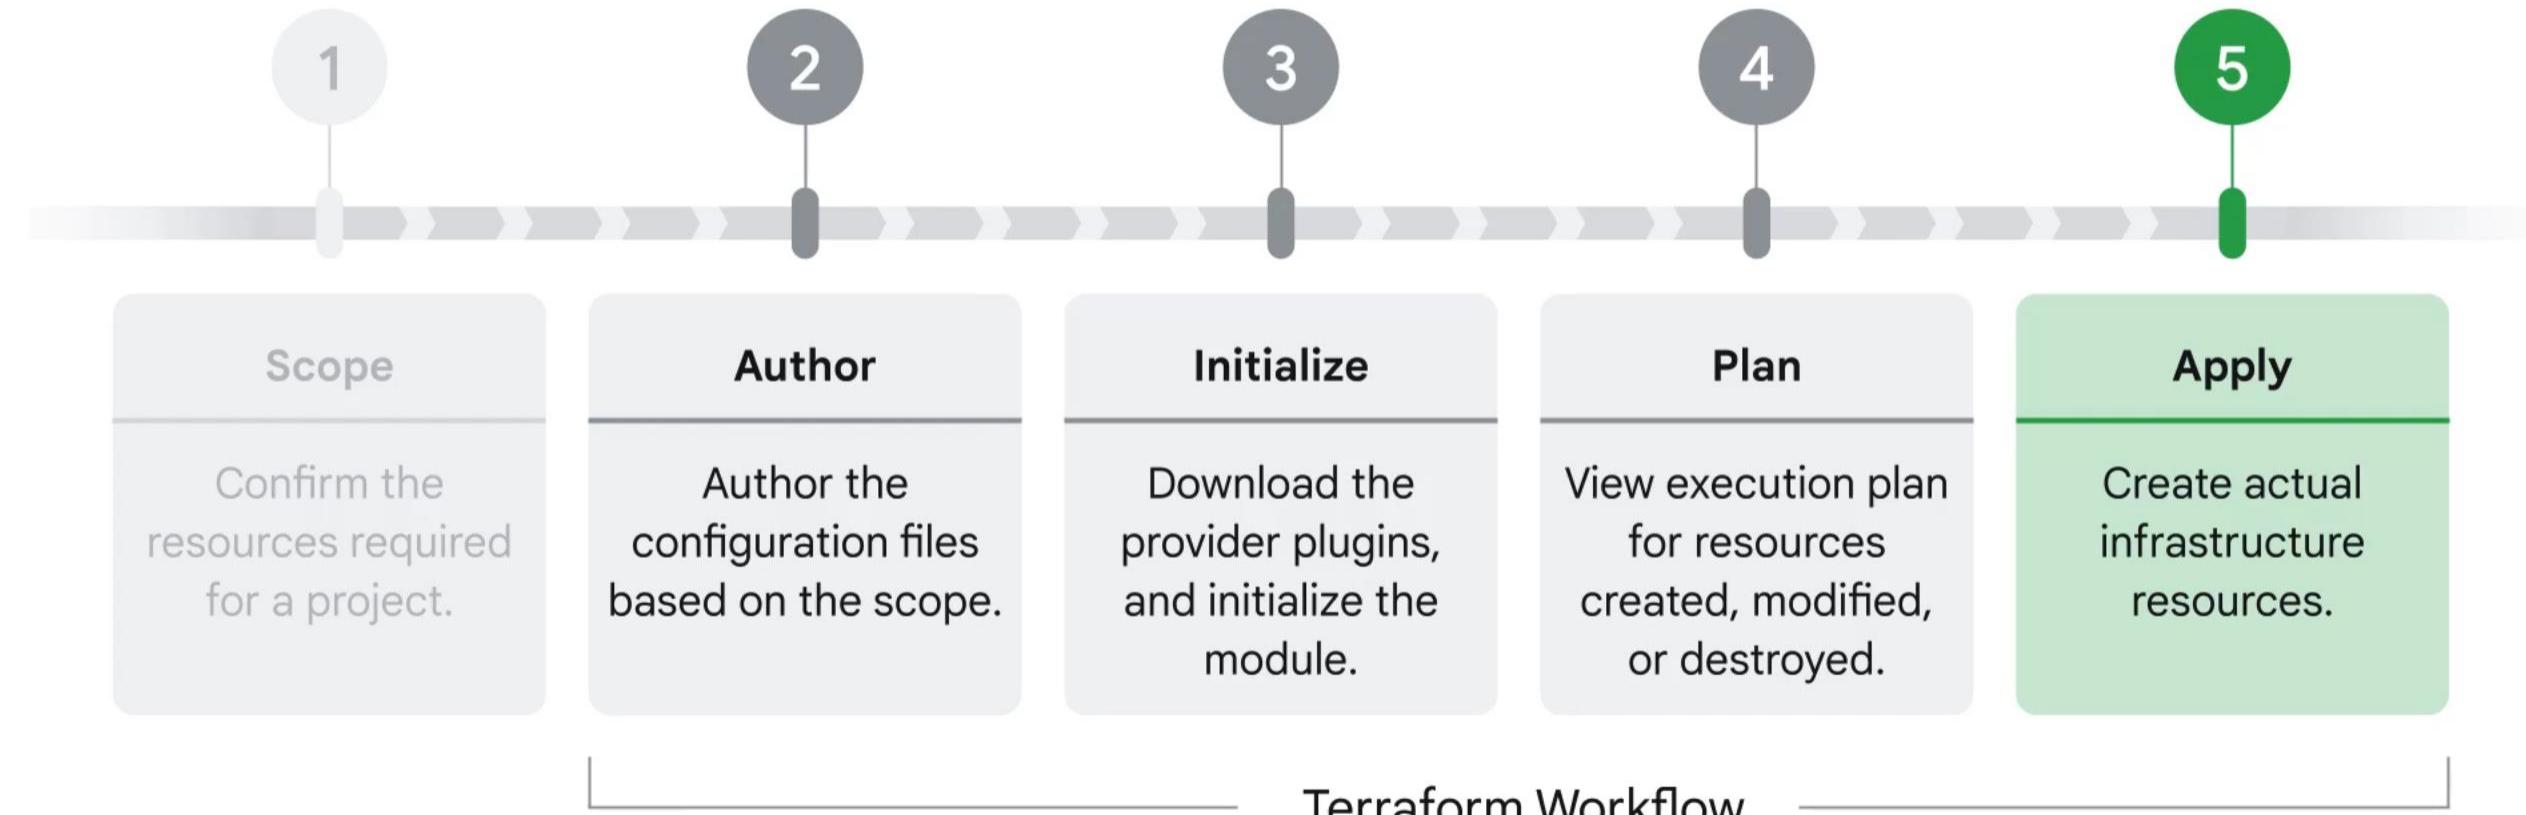
# Plan Phase:

`terraform plan` creates an execution plan

```
resource "google_storage_bucket" "example-bucket"{
  name      = "student0313ab04569a94"
  location  = "US"
}
```

```
$terraform plan
Terraform will perform the following actions:
# google_storage_bucket.example-bucket will be created
+ resource "google_storage_bucket" "example-bucket" {
    + force_destroy          = false
    + id                     = (known after apply)
    + location               = "US"
    + name                   = "student0313ab04569a94"
    + storage_class          = "STANDARD"
    + uniform_bucket_level_access = (known after apply)
}
Plan: 1 to add, 0 to change, 0 to destroy.
```

# Executes the actions proposed in a Terraform plan using **terraform apply**



# Apply Phase

`terraform apply` executes the plan

```
resource "google_storage_bucket" "example-bucket"{
  name      = "student0313ab04569a94"
  location  = "US"
}
```

```
$terraform apply
Terraform will perform the following actions:
# google_storage_bucket.example-bucket will be created
+ resource "google_storage_bucket" "example-bucket" {
    + force_destroy          = false
    + id                     = (known after apply)
    + location               = "US"
    + name                   = "student0313ab04569a94"
    ...
}

Apply changes: yes
google_storage_bucket.example-bucket: Creating...
google_storage_bucket.example-bucket: Creation complete after 1s [id=student0313ab04569a94]
Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```

# Code conventions

Formatting best practices:

- Separate meta arguments from the other arguments.
- Use two spaces for indentation.
- Align values at the equal sign.
- Place nested blocks below arguments.
- Separate blocks by one blank line.

```
resource "google_storge_bucket" "mybucket" {  
    boot_disk { #nested arguments above  
        initialize_params {  
            image = "debian-cloud/debian-9"  
        }  
    }  
    count = 2 #meta-argument in between  
    machine_type=e2_micro #unaligned equal signs  
    ..  
}  
  
resource "google_storge_bucket" "mybucket" {  
    count      = 2 #meta-argument last  
  
    machine_type  = e2_micro #align equal signs  
  
    boot_disk { #nested arguments below  
        initialize_params {  
            image = "debian-cloud/debian-9"  
        }  
    }..  
}
```



# terraform fmt

Terraform *fmt* can enforce code convention best practices.

Before terraform fmt:

```
resource "google_storge_bucket" "mybucket" {
    boot_disk { #nested arguments above
        initialize_params {
            image="debian-cloud/debian-9"
        }
    }
    count = 2 #meta-argument in between
    machine_type=e2_micro #unaligned equal signs
    ..
}
```

After terraform fmt:

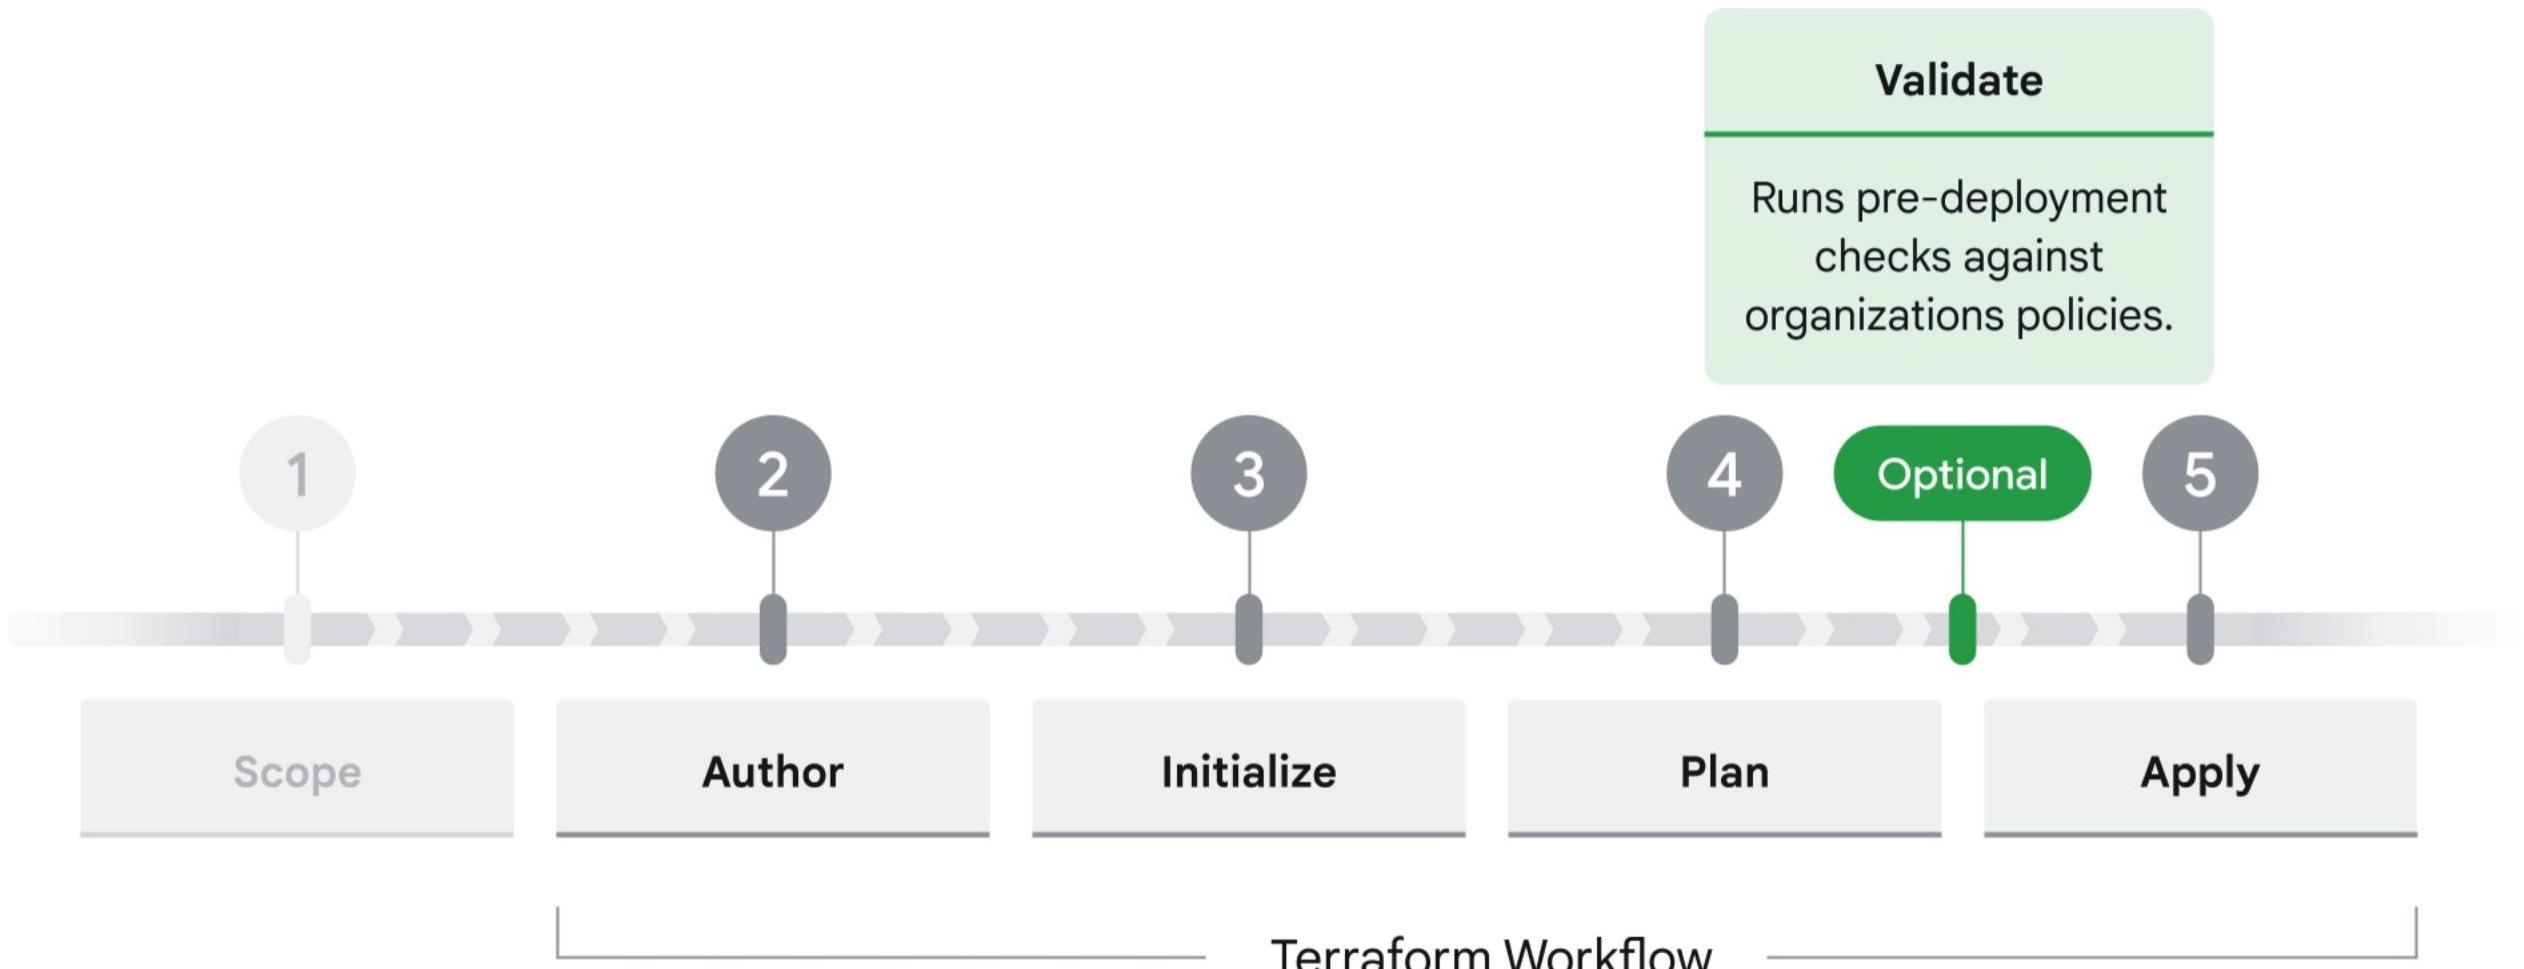
```
resource "google_storge_bucket" "mybucket" {
    count          = 2 #meta-argument first
    machine_type   = e2_micro #align equal signs
    #line space before a nested block

    boot_disk { #nested arguments below
        initialize_params {
            image = "debian-cloud/debian-9"
        }
    }
}
```

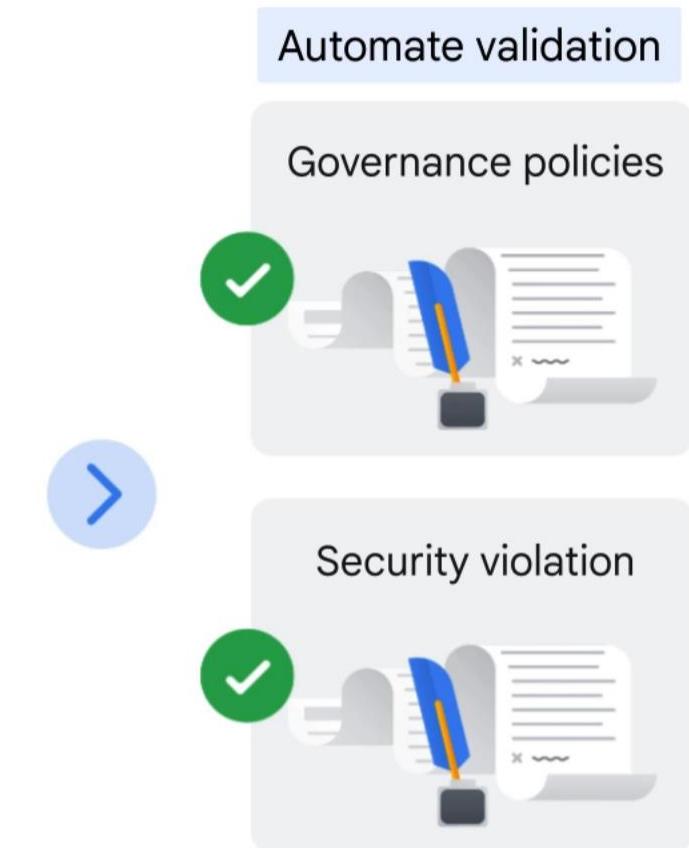
# terraform destroy command

```
$ terraform destroy
google_storage_bucket.example-bucket: Refreshing state... [id=student0313ab04569a94]
..
Terraform will perform the following actions:
# google_storage_bucket.example-bucket will be destroyed
- resource "google_storage_bucket" "example-bucket" {
    - force_destroy          = false -> null
    - id                     = "student0313ab04569a94" -> null
    - location               = "US" -> null
    - name                   = "student0313ab04569a94" -> null
    ..
}
Plan: 0 to add, 0 to change, 1 to destroy.
..
google_storage_bucket.example-bucket: Destroying... [id=student0313ab04569a94]
google_storage_bucket.example-bucket: Destruction complete after 0s
...
Destroy complete! Resources: 1 destroyed.
```

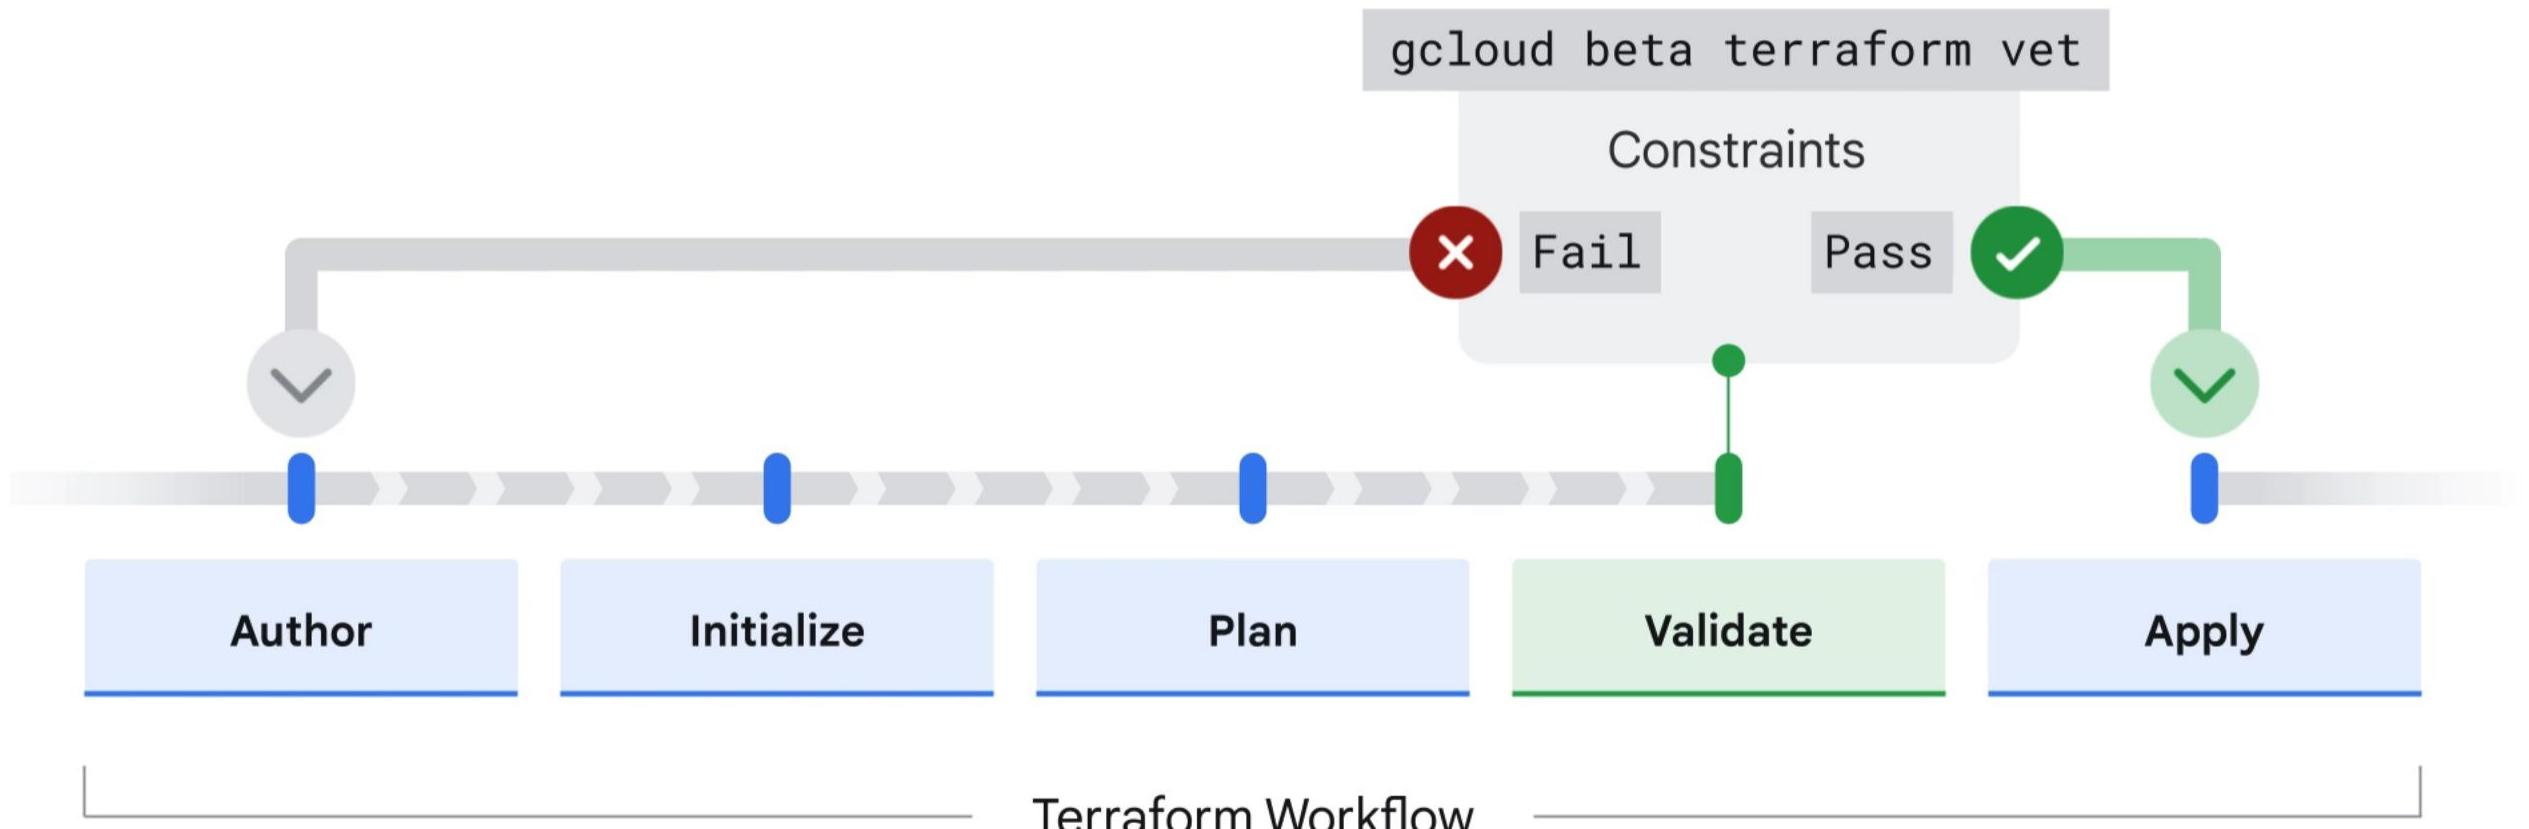
# The validate phase



# Why use the Terraform validator tool?



# The Terraform Validator



# Terraform Validator Benefits

Enforce policies	Enforce policies at any stage of application development
Remove manual errors	Remove manual errors by automating policy validation
Reduce time to learn	Reduce learning time by using a single paradigm for all policy management

# Terraform Validator uses

01

Platform teams can add guardrails to infrastructure CI/CD pipelines to ensure all changes are validated.

02

Application teams and developers can validate Terraform configuration with organization's central policy library.

03

Security teams can create a centralized policy library to identify and prevent policy violations.



# Potential Improvements



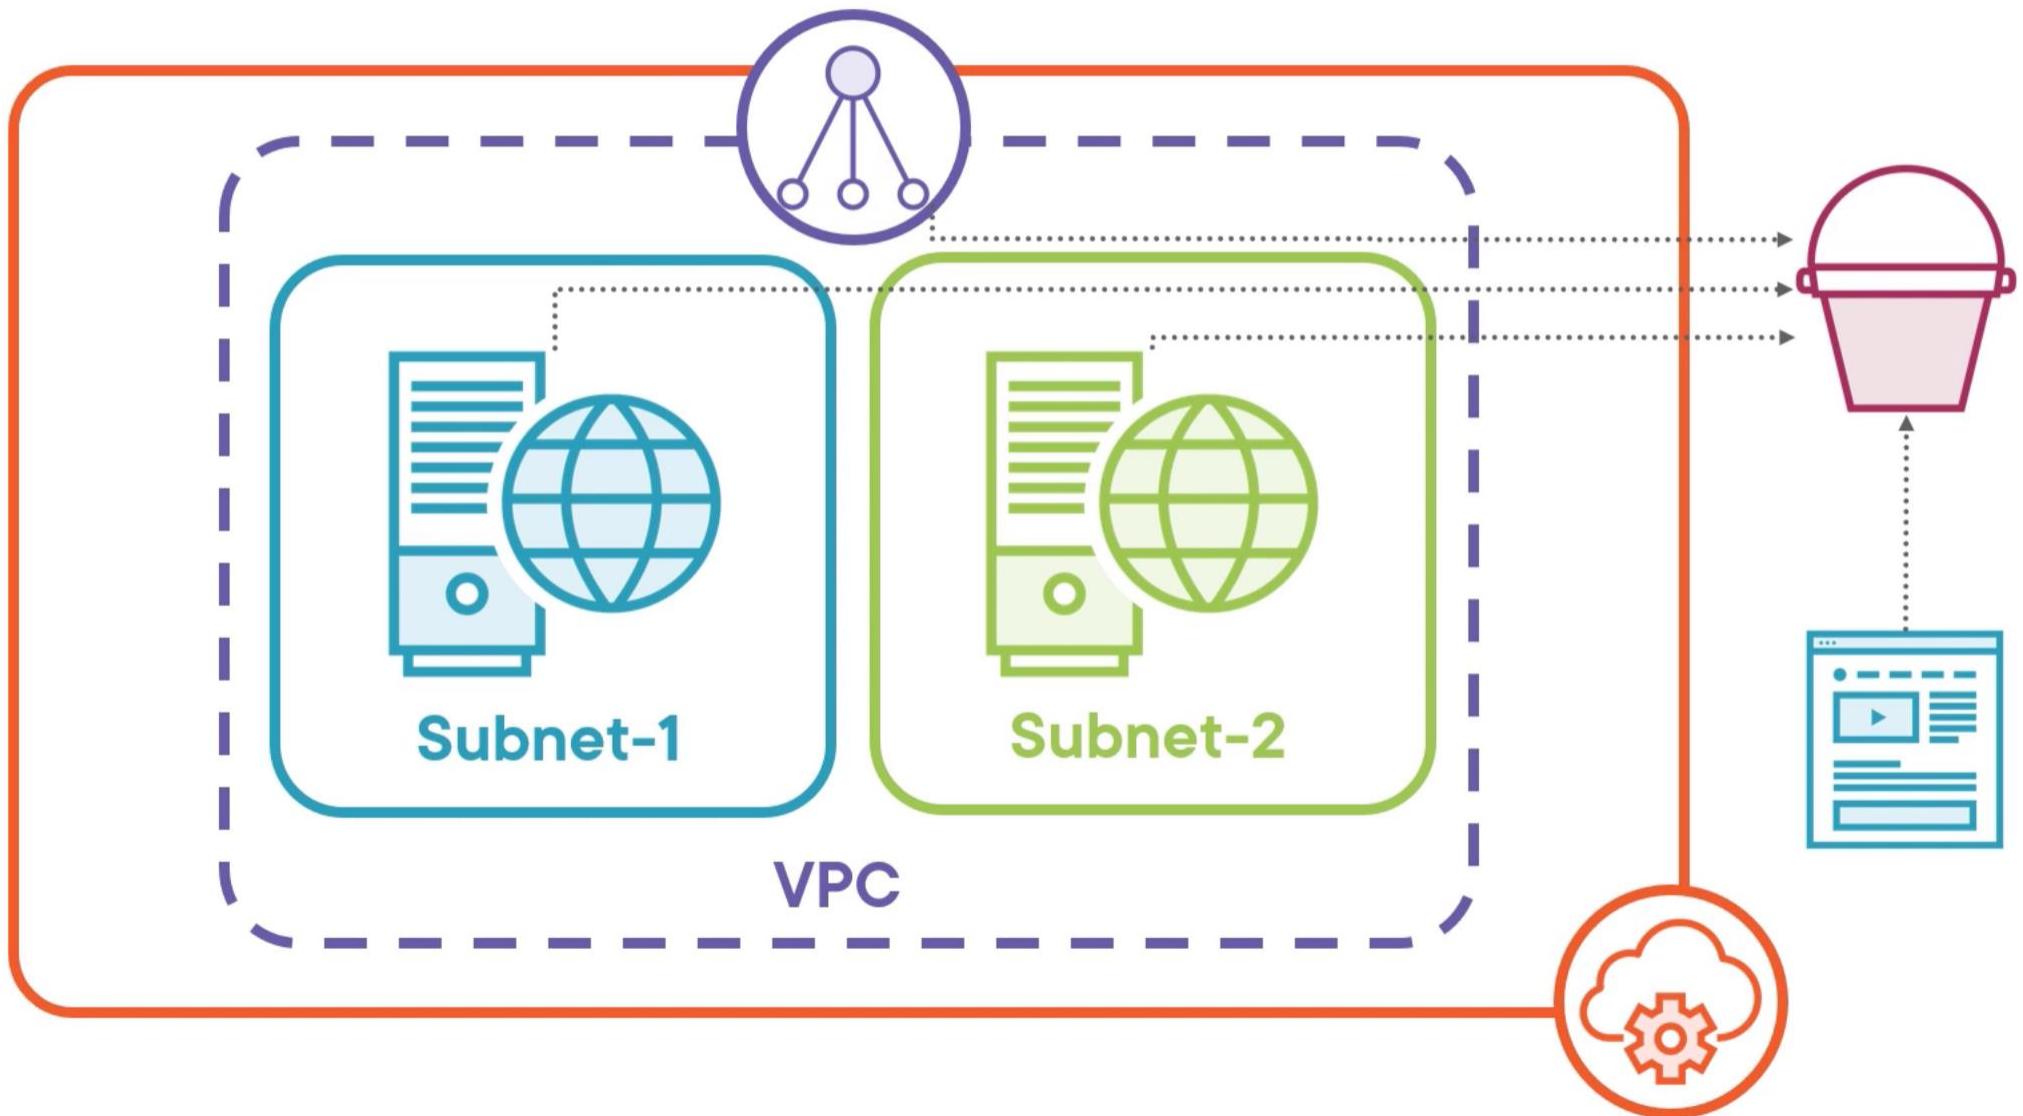
**Dynamically increase instances**

**Use a template for startup script**

**Simplify networking input**

**Add consistent naming prefix**

# Deployment Architecture



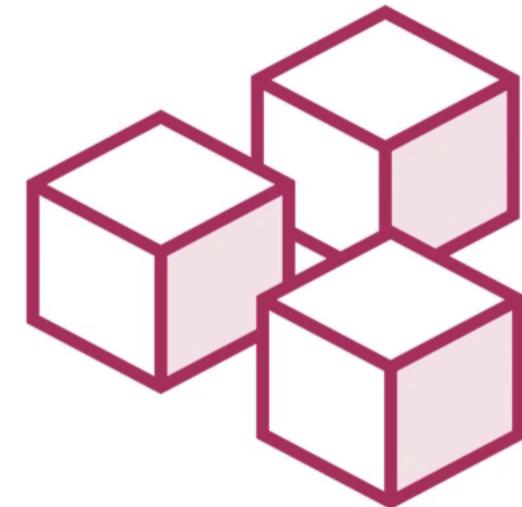
# Looping Constructs

[1,2,3]



Count  
Integer

For\_each  
Map or set



Dynamic blocks  
Map or set

```
resource "aws_instance" "web_servers" {  
  count = 3  
  tags = {  
    Name = "globo-web-${count.index}"  
  }  
}
```

## Count References

```
<resource_type>.<name_label>[element].<attribute>  
aws_instance.web_server[0].name # Single instance  
aws_instance.web_server[*].name # All instances
```

## For\_each Syntax

s3.tf

```
resource "aws_s3_bucket_object" "taco_toppings" {  
    for_each = {  
        cheese = "cheese.png"  
        lettuce = "lettuce.png"  
    }  
    key      = each.value  
    source   = ".${each.value}"  
    tags     = {  
        Name = each.key  
    }  
}
```

```
resource "aws_s3_bucket_object" "taco_toppings" {  
  for_each = {  
    cheese = "cheese.png"  
    lettuce = "lettuce.png"  
  }  
}
```

## For\_each References

```
<resource_type>.<name_label>[key].<attribute>  
aws_s3_bucket_object.taco_toppings["cheese"].id # Single instance  
aws_s3_bucket_object.taco_toppings[*].id # All instances
```

# Terraform Functions and Expressions

---

# Terraform Expressions



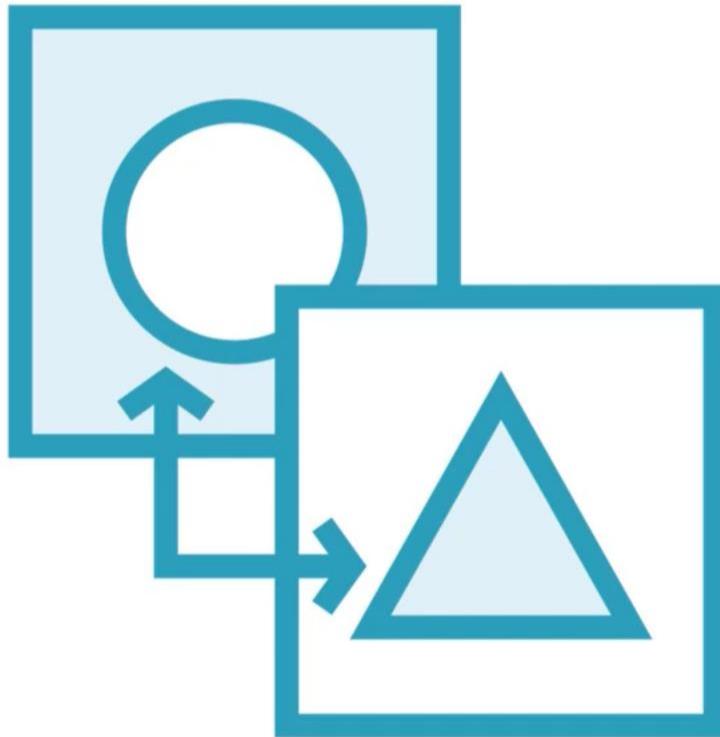
**Interpolation and heredoc**

**Arithmetic and logical operators**

**Conditional expressions**

**For expressions**

# Terraform Functions



**Built-in Terraform**

**Func\_name(arg1, arg2, arg3, ...)**

**Test in terraform console**

**Several broad categories**

# Common Function Categories

**Numeric**

`min(42, 13, 7)`

**String**

`lower("TACOS")`

**Collection**

`merge(map1, map2)`

**IP network**

`cidrsubnet()`

**Filesystem**

`file(path)`

**Type Conversion**

`toset()`

# Functions to Use

```
# Startup script  
  
templatefile(file_location, { map of variables })  
  
# Extract subnet address from VPC CIDR  
  
cidrsubnet(cidr_range, subnet bits to add, network number)  
  
# Add tags to common tags  
  
merge(common_tags, { map of additional tags })  
  
# S3 bucket name  
  
lower("bucket name")
```

# More Teams, More Problems



**Information Security**  
Define roles, policies,  
and groups

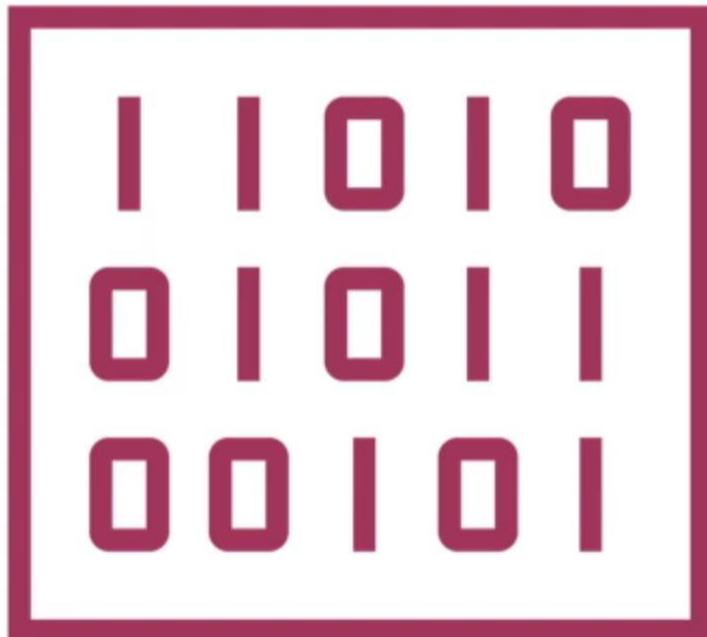


**Software Development**  
Read network  
configuration for app  
deployment



**Change Management**  
Store configuration  
data centrally

# Data Sources



**Glue for multiple configurations**

**Resources are data sources**

**Providers have data sources**

**Alternate data sources**

- Templates
- HTTP
- External
- Consul

# HTTP Data Source

```
# Example data source  
  
data "http" "my_ip" {  
    url = "http://ifconfig.me"  
}  
  
# Using the response  
data.http.my_ip.body
```

# Templates



**Manipulation of strings**

**Overloaded term**

- Quoted strings
- Heredoc syntax
- Provider
- Function

**Interpolation and directives**

```
# Simple interpolation
"${var.prefix}-app"

# Conditional directive
"%{ if var.prefix != "" }${var.prefix}-app%{ else }generic-app%{ endif }"

# Collection directive with heredoc
<<EOT
%{ for name in local.names }
${name}-app
%{ endfor }
EOT
```

## Template strings

You've been using these implicitly

Expressed in the configuration directly

Heredoc preferred for readability

## Template Syntax In-line

```
# Template data source
data "template_file" "example" {

    count = "2"

    template = "$${var1}-$$${current_count}"

    vars = {
        var1 = var.some_string
        current_count = count.index
    }
}

# Using the template
data.template_file.example.rendered
```

```
# Template configuration

data "template_file" "peer-role" {
    template = file("peer_policy.txt")
    vars = {
        vpc_arn = var.vpc_arn
    }
}

# Or templatefile function

templatefile("peer_policy.txt",
    { vpc_arn = var.vpc_arn } )
```

```
peer_policy.txt

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Action": [
                "ec2:AcceptVpcPeeringConnection",
                "ec2:DescribeVpcPeeringConnections"
            ],
            "Effect": "Allow",
            "Resource": [
                "${vpc_arn}"
            ]
        }
    ]
}
```

# Applying Policy as Code

---

# HashiCorp Sentinel

Policy as code

Version control

Native policy language

Fine grained and  
conditional

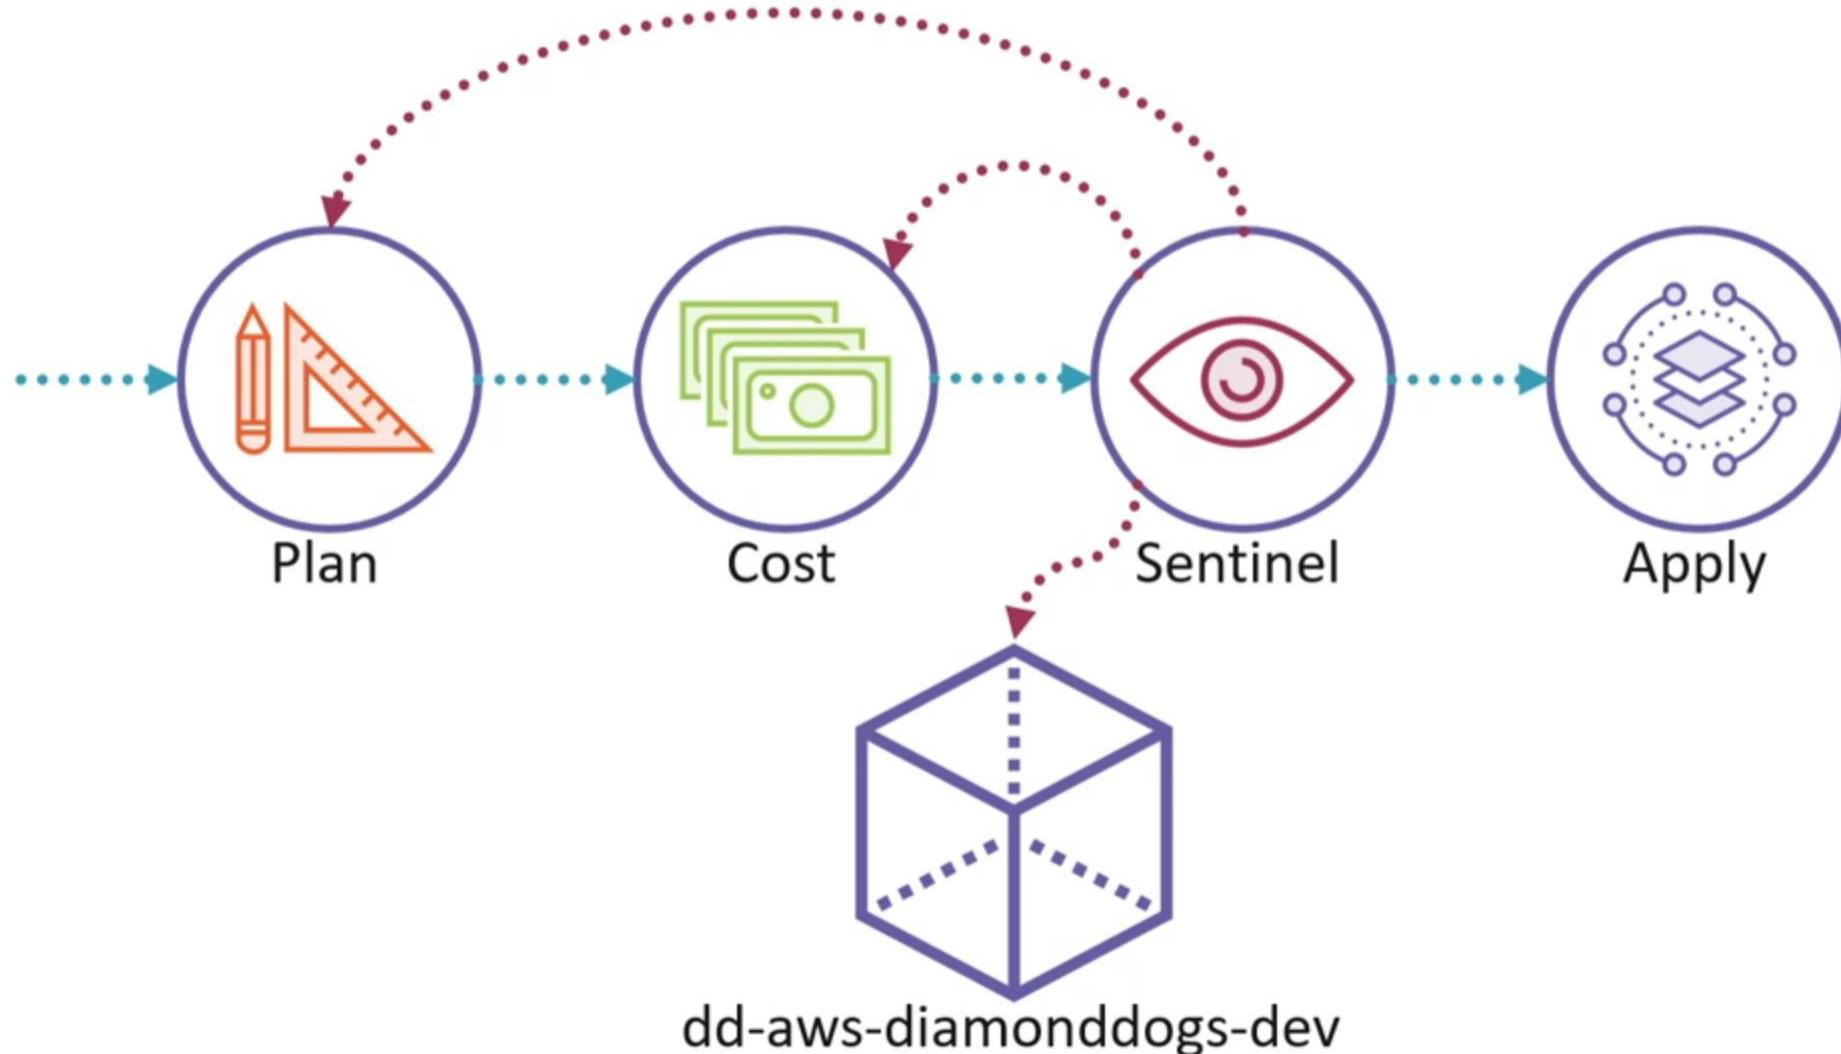
Import data sources

Enforcement levels

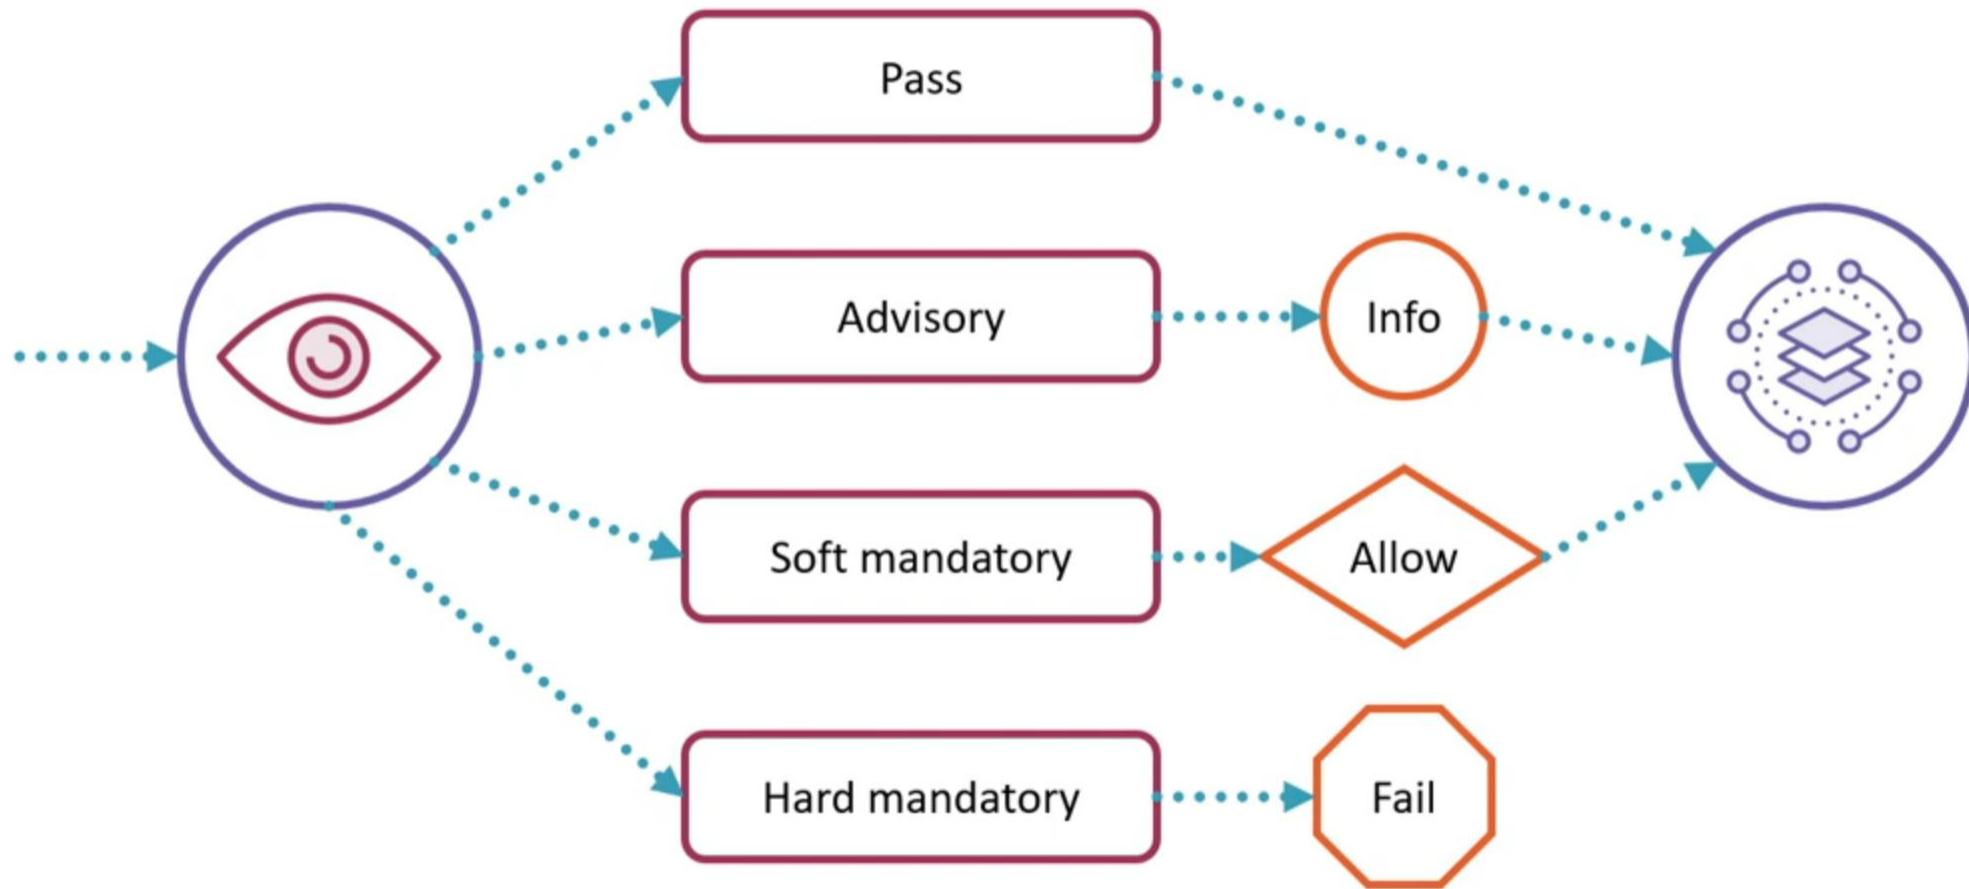
# Sentinel on Terraform Cloud



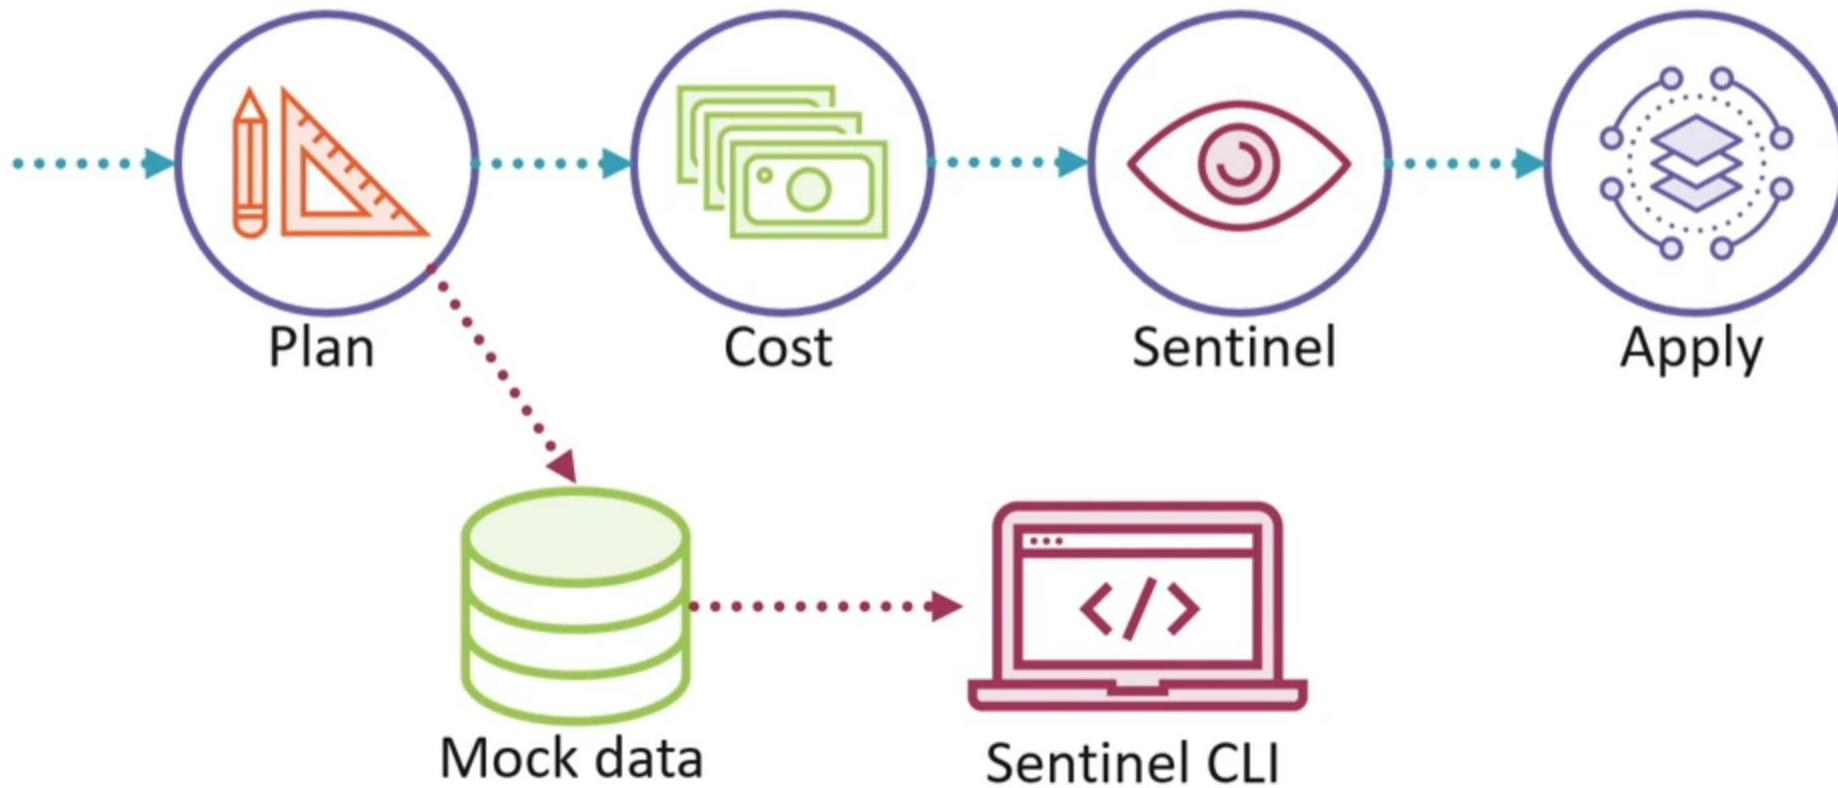
# Policy Evaluation



# Enforcement Actions



# Testing Sentinel



## instance-size.sentinel

```
# Import tfplan

import "tfplan/v2" as tfplan

# Get all EC2 instances

ec2_instances = filter tfplan.resource_changes as _, rc {
    rc.type is "aws_instance" and rc.mode is "managed" }

# Test the instance size

main = rule { all ec2_instances as _, instance {
    instance.change.after.instance_type in ["t2.micro", "t2.small"] } }
```

## sentinel.hcl

```
policy "instance-type" {  
    enforcement_level = "hard-mandatory"  
}  
  
policy "instance-tags" {  
    source = "https://raw.githubusercontent.com/.../instance-tags.sentinel"  
    enforcement_level = "advisory"  
}
```

# Using Modules

Leverage existing function libraries

**sentinel.hcl**

```
module "tfplan-functions" {  
  source =  
  "https://raw.githubusercontent.com/.../  
  tfplan-functions.sentinel"  
}
```

**instance-tags.sentinel**

```
import "tfplan-functions" as plan  
  
ec2_instances =  
  plan.find_resources("aws_instance")
```

---

# Questions ???